

Binary Classification in Unstructured Space With Hypergraph Case-Based Reasoning

Alexandre Quemy

IBM Krakow Software Lab, Cracow, Poland

Faculty of Computing, Poznań University of Technology, Poznań, Poland

Abstract

Binary classification is one of the most common problem in machine learning. It consists in predicting whether a given element is of a particular class. In this paper, a new algorithm for binary classification is proposed using a hypergraph representation. Each element to be classified is partitioned according to its interactions with the training set. For each class, the total support is calculated as a convex combination of the *evidence* strength of the element of the partition. The evidence measure is pre-computed using the hypergraph induced by the training set and iteratively adjusted through a training phase. It does not require structured information, each case being represented by a set of *agnostic information* atoms. Empirical validation demonstrates its high potential on a wide range of well-known datasets and the results are compared to the state-of-art. The time complexity is given and empirically validated. Its capacity to provide good performances without hyperparameter tuning compared to standard classification methods is studied. Finally, the limitation of the model space is discussed and some potential solutions proposed.

Keywords: binary classification, hypergraph, case-base reasoning

1. Introduction

In many real-life situations, one tries to take a decision based on previous *similar* situations. Each situation is described by a certain amount of

Email address: aquemy@pl.ibm.com (Alexandre Quemy)

information, either collected by an expert according to the relevance of the
5 information, or automatically by some sensors or algorithms. Those situations
share similarities on which to make analogies or counter-examples in order to
take a new decision. Conversely, in general, if two situations do not share any
common characteristic, then they are totally independent, i.e. it is impossible
to infer one’s outcome from the other one. The purpose of supervised machine
10 learning algorithms is to exploit the available information and interactions
between past cases or examples in order to build a model or infer the key
rules to take correct decisions.

Due to the large variety of concrete situations that can be reduced to binary
classification, it is one of the most studied problems in machine learning. In
15 this paper, we investigate the problem of binary prediction under a supervised
setting, i.e. given a history of previous situations labeled with the correct
output.

This paper **contributes** to binary classification with a new algorithm
called Hypergraph Case-Based Reasoning (HCBR). The idea is to create a
20 hypergraph where each element of the training set is a hyperedge and vertices
are represented by the features describing the elements. The intersections
between edges create a partition unique to a hypergraph. For each case, we
model the support as a convex combination of the elements of this partition.
Each of those elements is valued according to its importance w.r.t. to the set
25 of all the hyperedges it belongs to and their respective labels.

The paper is an extension of [1] with a focus on separating the abstract
framework from the specific implementation used for the experiments. On
top of the previous comparison to the best literature results, we extended the
empirical validation with a comparison between HCBR and 9 alternative
30 methods for binary classification, showing the robustness of our approach. A
study of learning curves and model spaces, highlighted some properties of
HCBR and the limitation of the current model space, thus shaping prior-
ity axes for future work. Among other refinements, we provided matricial
formulation for the model space and model selection to ease an efficient imple-
35 mentation. We also showed more evidence concerning the model calibration.

The well-known “no free lunch” theorem states that there is no unique
algorithm that can outperform all the others in supervised learning. The
advantages of an algorithm on a particular domain, problem or even instance
is “paid” on another. However, the case description and representation play
40 a preponderant role in the performances. Yet the algorithms are very often
constrained by being designed to a specific representation that may not be

suitable for a given problem or user data. The proposed method is totally agnostic about the data representation¹ and does not require the data to be structured. For instance wrongly, it can work on atomic representations
45 such as Bag-of-Word representations for texts where an atom is a single word (or n -gram). In this case, two elements would not have the same number of elements and it may be hard to properly define the feature domain.

The plan of the paper is as follows: in Section 2 the problem of binary classification is formalized and a more general formulation in an abstract space
50 of information is proposed. We illustrate the interest of such formulation by practical constraints. The main contribution of this paper is divided into two parts: Section 3 defines the general HCBR framework and its particular instance we used in practice. Section 4 presents empirical results on seven datasets from the UC Irvine Machine Learning Repository (UCI)² and the
55 LIBSVM³. The paper ends in Section 5 with a discussion about the results, future work, and possible improvements.

2. Binary Classification

Before introducing the problem of binary classification, we present the notations used throughout this paper. Vectors are denoted in bold and small
60 case (e.g. \mathbf{x}) and their component in small case (e.g. x_i). A collection of vectors is denoted in bold and large case (e.g. \mathbf{X}). $|\mathbf{x}|$ (resp. $|\mathbf{X}|$) denotes the cardinal of the vector \mathbf{x} (resp. the collection \mathbf{X}) while $\|\mathbf{x}\|$ is its norm. The scalar product is denoted by $\langle \cdot, \cdot \rangle$. For a matrix $A = (a_{ij})$, $a_{i\cdot}$ is the i th row vector and $a_{\cdot j}$ the j th column vector.

In machine learning, the problem of classification consists in finding a
65 mapping from an input vector space \mathcal{X} to a discrete decision space \mathcal{Y} using a set of examples. The binary classification problem is a special case such that \mathcal{Y} has only two elements. It is often viewed as an approximation problem s.t. we want to find an estimator \bar{J} of an unknown mapping J available only
70 through a sample called *training set*. A training set (\mathbf{X}, \mathbf{y}) consists of N input vectors $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ and their associated correct class $\mathbf{y} = \{J(\mathbf{x}_i)\}_{i=1}^N$.

Let $\mathcal{J}(\mathcal{X}, \mathcal{Y})$ be the class of mappings from \mathcal{X} to $\{-1, 1\}$, or simply \mathcal{J} if there is no ambiguity. A machine learning algorithm to solve binary

¹With the exception that the current version must work with discrete representation.

²<http://archive.ics.uci.edu/ml/index.php>

³<https://www.csie.ntu.edu.tw/~cjlin/libsvmtools/datasets/binary.html>

classification is an application $\mathcal{A} : \mathcal{X}^N \times \mathcal{Y}^N \mapsto \mathcal{J}$ capable of providing a
75 good approximation for any $J \in \mathcal{J}$ under some assumptions on the *quality* of
the training set. In practice, it is not reasonable to search directly in \mathcal{J} and
some assumptions on the “shape” of J are made s.t. $\bar{J} = \mathcal{A}(\mathbf{X}, \mathbf{y})$ belongs
to a *hypothesis* space or model space $\mathcal{H} \subset \mathcal{J}$. This restriction implies not
only that the exact mapping J is not always reachable but might also not
80 be approximated correctly by any element of \mathcal{H} . The choice of the model
space is thus crucial as it should be large enough to represent fairly complex
functions and small enough to easily find the best approximation of a given
mapping J .

In general, a robust classification algorithm must be able to approximate
correctly any possible mapping. The problem of finding such algorithm
consists in minimizing the generalization error for all possible mappings.
Formally, it consists in solving:

$$\min_{\mathcal{A}} \sum_{J \in \mathcal{J}} \int_{\mathcal{X}} \|J(\mathbf{x}) - \bar{J}(\mathbf{x})\| \mu(\mathbf{x}) dx \quad (1)$$

In practice, the generalization error cannot be computed and the set
of possible mappings is unreasonably large. For those reasons, we aim at
minimizing the empirical classification error on a reasonably large set of
datasets $\mathcal{D} = \{(\mathbf{X}_1, \mathbf{y}_1), \dots, (\mathbf{X}_K, \mathbf{y}_K)\}$, i.e.

$$\min_{\mathcal{A}} \sum_{(\mathbf{X}, \mathbf{y}) \in \mathcal{D}} \sum_{(\mathbf{x}, y) \in (\mathbf{X}, \mathbf{y})} \mathbb{I}_{\{y \neq \bar{J}(\mathbf{x})\}} \quad (2)$$

85 .
To highlight the differences between standard classification problem and
our approach, we briefly present linear models for classification followed by
the classification in unstructured space. We also present classification on
hypergraphs to discuss how the standard approach is not suitable for our
90 purposes.

2.1. Linear Binary Classification

The problem of binary classification is commonly studied with $\mathcal{X} = \mathbb{R}^M$.
Many popular classification approaches such as SVM, perceptron or logistic
regression define the model space as the set of M -hyperplan. A M -hyperplan

is uniquely defined by a vector $\mathbf{w} \in \mathbb{R}^M$ and a bias $w_0 \in \mathbb{R}$ and can be explicited by

$$h_{\mathbf{w}}(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle + w_0 = 0 \quad (3)$$

The *homogeneous* notation consists in adding w_0 to \mathbf{w} by rewritting \mathbf{x} such that $\mathbf{x} = (1, x_1, \dots, x_M)$. The hyperplan equation (3) is then expressed by $h_{\mathbf{w}}(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle$. A hyperplan separates \mathbb{R}^M into two regions, and thus, can be used as a discriminative rule such that

$$\bar{J}_{\mathbf{w}}(x) = \begin{cases} 1 & h_{\mathbf{w}}(\mathbf{x}) > 0 \\ -1 & h_{\mathbf{w}}(\mathbf{x}) \leq 0 \end{cases} \quad (4)$$

Then, given a training set (\mathbf{X}, \mathbf{y}) , the classification problem is equivalent to find the best hyperplan such that it minimizes a certain *loss* function over the training set:

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \sum_{i=1}^N L(\mathbf{w}, \mathbf{x}_i) + \lambda R(\mathbf{w}) \quad (5)$$

where $R(\mathbf{w})$ is *regularization* term to prevent overfitting (usually $\|\mathbf{w}\|_2^2$ or $\|\mathbf{w}\|_1$) and $\lambda > 0$ a hyperparameter controlling the effect of regularization. Note that (5) is a parametric problem due to the choice of model space. Several losses functions exists, and are generally based on the *margin* of \mathbf{x}_i which is defined by $m(\mathbf{w}, \mathbf{x}_i) = J(\mathbf{x}_i)h_{\mathbf{w}}(\mathbf{x}_i)$. The margin represents the distance of a vector \mathbf{x}_i to the hyperplan defined by $h_{\mathbf{w}}$. It is positive if \mathbf{x}_i is correctly classified, negative otherwise. For the most known, the 0-1 loss, hinge loss and log loss are defined by

$$\begin{aligned} L_{01}(\mathbf{w}, \mathbf{x}) &= \mathbb{I}_{\{m(\mathbf{w}, \mathbf{x}) \leq 0\}} \\ L_{\text{hinge}}(\mathbf{w}, \mathbf{x}) &= \max(0, 1 - m(\mathbf{w}, \mathbf{x})) \\ L_{\text{log}}(\mathbf{w}, \mathbf{x}) &= \ln(1 + e^{-m(\mathbf{w}, \mathbf{x})}) \end{aligned} \quad (6)$$

The Perceptron algorithm uses the 0-1 loss, while SVM minimizes the hinge loss and the Logistic Regression the log loss.

2.2. On the necessity of unstructured spaces

95 In this paper, we relax the three following implicit constraints on the input vector space:

1. we do not assume to have any inner product or metric embedded with the input vector space,
- 100 2. we do not assume the cardinality of the input vectors, such that it is possible to build a model and make predictions on incomplete systems (for instance missing values in some rows in a database or Bag-of-Words representation),
3. we do not assume anything about the concrete *representation* of features while in most classification methods, all the features belongs to the same space, often \mathbb{R} .

Most of classification methods work directly in the feature space, or, when they modify it, the resulting space inherits the topology of the initial space. In the particular case of linear models, a scalar product is mandatory to represent the decision itself and to quantify the distance from an element to the decision boundary. In some cases, a metric is obviously available, and in some other cases some ad-hoc metrics may be constructed for a specific problem. However, there a plenty of practical situations where the *good* metric is not known (e.g. it requires a very specific expertise). Also, it is very common to miss some data in a dataset (e.g. due to a defective sensor for some period of time), or to have to include new pieces of information after an initial model was built (e.g. adding new sensors). In the case of SVM, for instance in \mathbb{R}^3 , if an element has only 2 components it does not describe a point but a plan and as a result, there is no way separate it with a (hyper)plan.

120 Conversely, consider a model built in \mathbb{R}^2 with some additional information available after such that the classification instance now lives in \mathbb{R}^3 . Mathematically, there is no guarantee that the model in two dimensions is even close to the model built from scratch in three dimensions. For instance, if the points are not linearly separable in two dimensions but are in three dimensions, there is no chance for the projection of the plane into the subspace of two dimensions to be the model found while working only in this subspace. Despite multiple algorithms have an *online* counterpart, i.e. the model keeps being trained with new input vectors, as far as we know, there is no significant work on having this *horizontal* scalability (as opposed to vertical scalability).

130 One step toward horizontal scalability is to work only in unstructured spaces. Last, the features representation is very often imposed by the classification method itself and when the dataset does not respect those requirements, it has to be transformed. Of course, exploiting the specificities of features often

lead to better results but also requires manual expertise that blocks a wider
 135 adoption by end-users. This is particularly true with categorical variables
 that often require to be encoded taking into account the method specificity.
 For instance, for Decision Tree and its Ensemble counterpart, the Random
 Forest, the accuracy depends on the encoder and using a Numerical encoder,
 a One-Hot encoder or a Binary encoder may drastically change the model
 140 fitted on the same data. For this reason, it is also important to develop
 methods that are “agnostic” to the data representation or at least as less
 sensitive as possible.

2.3. Binary Classification In Unstructured Spaces

With all those considerations in mind, we reformulate the problem of
 binary classification for unstructured space. The idea is that the only *ad-hoc*
 operation we allow is checking if a particular feature belong to the element to
 classify, and by extension we allow elementary set operations. We consider a
 countable space \mathbb{F} and define $\mathcal{X} = 2^{\mathbb{F}}$. By abuse of notation, we call $\mathbf{x} \in \mathcal{X}$
 an “input vector” despite \mathcal{X} is not a vector space. We also refer to \mathbf{x} as a
case. In practice, it is very likely that only a subset of $2^{\mathbb{F}}$ may appear (for
 instance if two features encode two contradictory propositions or if every case
 has the same number of features). The real class for any input vector \mathbf{x} of $2^{\mathbb{F}}$
 is given by the mapping:

$$J: 2^{\mathbb{F}} \rightarrow \{-1, 1\} \quad (7)$$

$$\mathbf{x} \mapsto J(\mathbf{x}) \quad (8)$$

Assuming the unknown mapping takes value from the powerset of \mathbb{F} allows us
 145 not to have to know \mathbb{F} at all. We could propose a more general formulation,
 and in particular to handle continuous space, s.t. \mathcal{X} is defined as a σ -algebra
 on \mathbb{F} . However, this paper focuses on the specific case where \mathbb{F} is supposed
 countable and the σ -algebra is defined by the powerset of \mathbb{F} . We let the
 general case for future work.

150 Notice that in this paper we do not consider *uncertainty*: if two situations
 are described the same in \mathbb{F} , then they have the same label.

2.4. Classification and Hypergraph

Hypergraphs generalize graphs and can be used to represent higher order
 relations while graphs are limited to binary relations. A hypergraph is defined
 155 by a set of vertices and a collection of hyperedges where each hyperedge

is a subset of this set of vertices. Therefore, a graph is a special case of hypergraph for which each hyperedge contains only two vertices. We will formally introduce hypergraphs in Section 3.1. Recently hypergraphs have been used as data representation, and some classification algorithms on hypergraph have been proposed. A vast majority of approaches models the objects to classify as the set of vertices and constructs the hyperedges as the representation of a metric. This conventional approach is known as *neighborhood-based* hypergraph. The metric relies on some assumptions on the data or a specific representation (e.g. Zernike moment and histogram of oriented gradient to measure the distance between images in [2]) and for each vertex, a hyperedge is created to represent its k -nearest neighbors [3].

The problem of classification on hypergraph consists in labeling some unlabeled vertices given a training set such that all vertices in a same hyperedge have the same label. As all the vertices are known a priori, the problem is part of transductive learning. To learn the labels, the standard approach is to minimize a cost function based on a hypergraph equivalent of a graph Laplacian [2, 4] with a structural risk:

$$C(\mathbf{x}) = \mathbf{x}^t \Delta \mathbf{x} + \mu \|\mathbf{x} - \mathbf{y}\|^2 \quad (9)$$

where Δ is the hypergraph Laplacian, $\mu > 0$ a regularization factor and $\|\cdot\|$ a norm. The vector \mathbf{y} represents the initial labels for all vertices with $y_i = 0$ for unlabeled vertices, a negative (resp. positive) value for label -1 or 1.

On the contrary, the **method proposed in this paper** models the elements to classify as the hyperedges and the vertices as the different components of those elements. As far as we know, there is no previous work that uses this modeling choice. In addition, it does not require knowing all the elements before building the model: our approach is inductive. Contrary to the literature that focuses on the method to build the hyperedges, the proposed framework has a straightforward hypergraph construction and rather focus on model selection. Even more important, as most previous work consists in building metrics based on the features representation, it obviously conflicts with our goal of agnosticity described in the previous section.

3. Hypergraph Case-Based Reasoning

In this section, we introduce our main contribution with a **new framework for binary classification in unstructured spaces** called Hypergraph Case-Based Reasoning (HCBR). The presentation is broken down into

five steps. First, we present how to represent our training set as a hypergraph
 185 and how to project any element of $2^{\mathbb{F}}$ onto it. In a second part, we formally
 introduce the model space, the model parameters and the decision rule. Sec-
 tion 3.3 is dedicated to the parameter estimation while Section 3.4 focuses
 on the training and decision rule refinement. Before presenting experimental
 results, Section 3.5 provides the time complexity of the three main phases of
 190 the algorithm.

3.1. Representation And Projection

Before defining the projection operator used by HCBR to make prediction,
 we recall the definitions of a hypergraph and induced hypergraph as found
 in [5]. For additional results on hypergraphs, we refer the reader to this
 195 reference.

Definition 3.1 (Hypergraph). A hypergraph is defined by $H = (V, \mathbf{X})$ with
 V a set of vertices, \mathbf{X} the hyperedges such that $\forall \mathbf{x} \in \mathbf{X}, \mathbf{x} \subseteq V$.

A hypergraph can be viewed as a collection of subsets \mathbf{X} of a given set of
 vertices V . It is sometimes convenient to define a hypergraph solely by a
 200 collection of sets. In this case, the set of vertices, denoted $V_{\mathbf{X}}$, is implicitly
 defined as the union of edges.

Definition 3.2 (Induced Subhypergraph). The subhypergraph $H[A]$ induced
 by $A \subseteq V$ is defined by $H[A] = (A, \mathbf{X}_A)$ with $\mathbf{X}_A = \{\mathbf{x} \cap A \mid \mathbf{x} \cap A \neq \emptyset\}$.




Figure 1: The family $\mathcal{E} = \{\mathbf{e}_i\}_i$ forms the partition obtained by the union of the projection
 of cases and represents how the three cases share information.

A set of examples \mathbf{X} can be seen as a hypergraph $H = (\mathbb{F}, \mathbf{X})$, i.e. such
 205 that each example is a hyperedge (Figure 1). In practice, we do not need to

know \mathbb{F} as we can always use the implicit hypergraph $H = (\mathbb{F}_{\mathbf{X}}, \mathbf{X})$ where $\mathbb{F}_{\mathbf{X}}$ is the restriction of \mathbb{F} to the features that appear in the sample \mathbf{X} . For any hypergraph $H = (\mathbb{F}_{\mathbf{X}}, \mathbf{X})$, there exists a unique partition $\mathcal{E}_H = \{\mathbf{e}_i\}_{i=1}^M$ defined by the intersections of its edges as illustrated by Figure 1.

210 The idea of the projection of a case \mathbf{x} over an hypergraph is to return the partition of \mathbf{x} defined by the intersections with \mathcal{E}_H . In particular, the projection of a hyperedge is included in \mathcal{E}_H .

Definition 3.3 (Projection defined by a hypergraph). The projection operator π_H over a hypergraph $H = (V, \mathbf{X})$ for any $A \subseteq V$ is defined by
 215 $\pi_H(A) = \{\mathbf{e} \subseteq A \mid \exists \mathbf{X}' \subseteq \mathbf{X}_A, \mathbf{e} = \bigcap_{\mathbf{x} \in \mathbf{X}'} \mathbf{x}\}.$

Example: Each element of $\pi_H(\mathbf{x})$ is a (sub)set of features. For instance, in Figure 1, $\pi_H(\mathbf{x}_1) = \{\mathbf{e}_1, \mathbf{e}_2, \mathbf{e}_3, \mathbf{e}_6\}$ and in Figure 2, the projection of \mathbf{x} (in yellow) on H is represented by the family $\{\mathbf{e}'_i\}_i$.

220 We call *discretionary features* of \mathbf{x} (w.r.t. H) the (possibly empty) set of features that are not in $V_{\mathbf{X}}$, denoted $D_{\mathbf{x}}$. It can be interpreted as the features of \mathbf{x} that do not belong to any hyperedge. If a hypergraph induced by a set of examples represents a knowledge base at a given moment, the discretionary features of \mathbf{x} are new pieces of information that were never encountered before.

225 **Example:** Considering the hypergraph composed of \mathbf{x}_1 and \mathbf{x}_2 as illustrated by Figure 1, the set of discretionary features of \mathbf{x}_3 is \mathbf{e}_4 . In Figure 2, the yellow case \mathbf{x} has no discretionary feature: all its features are present at least in one example.

230 For any set of features $\mathbf{x} \subseteq \mathbb{F}$, we define $d_H(\mathbf{x}) = \{\mathbf{x}' \in X \mid \mathbf{x} \cap \mathbf{x}' \neq \emptyset\}$ the set of edges sharing some features with \mathbf{x} . In particular, the set d_H can be split into $d_H^{(1)}$ and $d_H^{(-1)}$ depending on the label of the case and defined by $d_H^{(l)}(\mathbf{x}) = \{\mathbf{x}' \in d_H(\mathbf{x}) \mid J(\mathbf{x}') = l\}$. $|d_H(\mathbf{x})|$ corresponds to the number
 235 of cases the case \mathbf{x} intersects while $|d_H^{(l)}(\mathbf{x})|$ counts the number of times the class l is used among this set of intersecting cases. Note that if $\mathbf{x} \notin \mathbf{X}$ and $|d_H(\mathbf{x})| = 0$, then $\mathbf{x} = D_{\mathbf{x}}$, i.e. the case \mathbf{x} is in relation with no case in \mathbf{X} . In the hypergraph literature, $|d(\mathbf{x})|$ is called the *degree* of a hyperedge and its domain of definition is restricted to \mathbf{X} . Finally, to use matricial notation, we
 240 define the vectors $\mathbf{d}^l = (|d^{(l)}(\mathbf{x}_i)|)_i^N$.




Figure 2: The projection of \mathbf{x} on H is represented by the family $\{\mathbf{e}'_i\}_i$. Under the projection lies a graph representation of \mathbf{x} with the partition elements $\{\mathbf{e}_i\}_i$ and their respective connections to the cases $\{\mathbf{x}_i\}_i$, in particular, $D_{\mathbf{x}} = \emptyset$.

From now, we consider only the specific hypergraph generated by the set of examples \mathbf{X} . For the sake of readability, we remove the subscript H .

3.2. Model Space

As discussed in Section 2.2, we relaxed some implicit constraints on the input vector space. As a counterpart, HCBBR relies on two assumptions: (i) the correct class of an input vector is the result of its features only and (ii) if two input vectors \mathbf{x} and \mathbf{x}' do not share any feature, they are *independent* i.e. \mathbf{x} cannot help to understand the correct class of \mathbf{x}' and vice versa. If (i) is a very common assumption, often implicit, (ii) comes from the fact there is no metric on \mathbb{F} to determine a distance between two elements s.t. we can rely only on intersections. As a result, HCBBR produces only *local* models because if a new input vector is independent of all examples, it is impossible to generate a prediction. On the contrary, a hyperplan model is *global* in a sense that it can produce a prediction for any element of \mathbb{R}^M . An example of concrete situation for which such assumptions are natural is a justice trial. Cases are composed of some elements, and the correct label is the result of a reasoning that can possibly use analogies or counter-examples with a set of past cases on top of the legal texts. However, if a judge or lawyer wants to use \mathbf{x} to justify the outcome of \mathbf{x}' , \mathbf{x}' must have similarities with \mathbf{x} .

This leads to define a model space such that the support toward one class for a given case \mathbf{x} is only function of its projection on the hypergraph. First, we give the **intuition** behind our model space. For each element \mathbf{e}_i in

the projection, we value 1) its importance in the case \mathbf{x} and 2) its support toward a specific class given the whole hypergraph. Intuitively, 1) consists in answering the question: *how important is \mathbf{e}_i w.r.t. \mathbf{x} ? or what is the potential for analogies or counter-examples between \mathbf{x} and the other examples also containing \mathbf{e}_i ?* The purpose of 2) is to value how important is \mathbf{e}_i with regards to the other elements of \mathcal{E} . For instance, to explain the case \mathbf{x} in Figure 1, we will use only the elements of the projection $\pi_H(\mathbf{x}_1) = \{\mathbf{e}_1, \mathbf{e}_2, \mathbf{e}_3, \mathbf{e}_6\}$. Without prior information, we will use the size of \mathbf{e}_i in \mathbf{x} to measure their importance \mathbf{x} : \mathbf{e}'_6 is the most important group of features and \mathbf{e}'_2 the least important. The way of measuring the support with regard to the whole hypergraph as well as an intuitive interpretation will be given in Section 3.3.

Let us now **formally** define the model space. Given the hypergraph $H = (\mathbb{F}, \mathbf{X})$ defined by a training set \mathbf{X} and its associated partition $\mathcal{E} = \{\mathbf{e}_i\}_i^m$ obtained using the projection defined in the previous section, the relation between an example \mathbf{x} and its class is modeled by

$$\begin{cases} s_{w,\mu}(\mathbf{x}) &= \sum_{i=1}^M w(\mathbf{e}_i, \mathbf{x}) \mu(\mathbf{e}_i, \mathbf{X}, \mathbf{y}) \\ \sum_{i=1}^M w(\mathbf{e}_i, \mathbf{x}_j) &= 1 \quad \forall 1 \leq j \leq N \\ \sum_{i=1}^M \mu(\mathbf{e}_i, \mathbf{X}, \mathbf{y}) &= 1 \end{cases} \quad (10)$$

where $w(\mathbf{e}_i, \mathbf{x}_j) \geq 0$ model the importance of \mathbf{e}_i in \mathbf{x}_j and $\mu(\mathbf{e}_i, \mathbf{X}, \mathbf{y})$ the support of \mathbf{e}_i for class 1 w.r.t. whole training set. For this reason, we call μ the *intrinsic strength* of \mathbf{e}_i . The assumption (ii) implies that if $\mathbf{e}_i \cap \mathbf{x}$ then $w(\mathbf{e}_i, \mathbf{x}) = 0$. For readability, we write s in place of $s_{w,\mu}$.

The classification rule consists in selecting the class with the total highest support:

$$\bar{J}_{w,\mu}(x) = \begin{cases} 1 & s_{w,\mu}(\mathbf{x}) > 0 \\ -1 & s_{w,\mu}(\mathbf{x}) \leq 0 \end{cases} \quad (\text{R1})$$

The strength function $s_{w,\mu}$ models the support toward class 1 deduced from the training set. The class of \mathbf{x} is 1 if and only if the convex combination of the strength is higher than 0. The classification problem consists then in finding the couple (w, μ) that minimizes the classification error:

$$(w^*, \mu^*) = \arg \min_{w, \mu} \sum_{(\mathbf{x}, \mathbf{y}) \in (\mathbf{X}, \mathbf{Y})} \mathbb{I}_{\{y \neq \bar{J}(\mathbf{x})\}} \quad (11)$$

The problem described by (11) is a functional problem, in general much harder than parametric ones such as (5), and without restrictions on the search space it may not be reasonable to look for a solution. For this reason and for
285 the rest of this paper, we will fix w a priori. As we do not formulate any particular assumption on the feature space and allow only basic set operations, a natural yet trivial way of modelling the importance of \mathbf{e}_i in \mathbf{x}_j is to set $w(\mathbf{e}_i, \mathbf{x}_j) = \frac{|\mathbf{x}_j \cap \mathbf{e}_i|}{|\mathbf{x}_j|}$. However, if one has some information about the impact of some features in some particular cases, w might be redefined to integrate this
290 prior knowledge.

Once the method to calculate w and μ is specified, and for a specific training set (\mathbf{X}, \mathbf{y}) , the model can be expressed matricially by

$$\begin{cases} S &= W\mu & W \in \mathcal{M}(\mathbb{R})_{N \times M}, \mu \in \mathbb{R}^M \\ ||\mathbf{w}_{j:}||_1 &= 1 & \forall 1 \leq j \leq N \\ ||\mu||_1 &= 1 \end{cases} \quad (12)$$

In addition, to fit a model capable of generalization, it is not enough to simply minimize the classification error. Assuming W is fixed and S set to the expected correct label (or even an arbitrary value such that each element is correctly classified), μ can be obtained by W^+S where W^+ is the
295 Moore-Penrose inverse⁴. But by doing this, S and μ holds no discriminative information and there is no chance it can provide good results on new cases. For this reason, μ has to be designed such that S is *calibrated*, i.e. the higher S is, the more confident the model is about the prediction. Ideally, among all the elements with e.g. a (normalized) strength close to 0.8, 80% of them
300 should be correctly classified.

To summarize, after defining a general model space to solve the binary classification problem in unstructured space, we made additional assumptions to reduce the problem of model selection to find μ such that the classification rule (R1) provides a good solution to the original problem (1). For this
305 purpose, we will proceed in three steps:

1. Define μ to capture as much information as possible from \mathcal{E} for any \mathbf{X} (Section 3.3).
2. Train the model to adjust the intrinsic strength on a specific X using the classification rule (R1) (Section 3.4).

⁴Depending on S and W it might be impossible to classify correctly all the elements using $\mu = W^+S$, however it probably returns a near-optimal solution.

- 310 3. Refine the classification rule (R1) to take into account the local nature
of the model (Section 3.4).

A summary and high level view of HCBR is given by Algorithm 1.

Algorithm 1 HCBR (High level view)

- 1: Build H and \mathcal{E} from \mathbf{X} .
 - 2: Calculate w and μ on \mathcal{E} .
 - 3: Adjust μ with training algorithm 2 on \mathbf{X} using rule (R1)
 - 4: **for** each \mathbf{x} in the test set **do**
 - 5: Calculate the projection $\pi(\mathbf{x})$.
 - 6: Calculate the support $s(\mathbf{x})$ using the projection.
 - 7: Predict using the updated rule (R2).
 - 8: **end for**
-

While SVM aims at separating the data using a single hyperplan in the original input vector space, HCBR tries to explain the decision for each case
315 by a convex combination expressed in a lower dimensional space \mathcal{E} generated by the data. In addition, the convex combinations depend on each other since the elements of \mathcal{E} are by definition the features resulting in the case intersections. For any case, the decision rule is a combination of the strength of the elements of its projection on the hypergraph.

320 *3.3. Model Selection*

In this section, we define how to concretely select μ for a given hypergraph. Our approach is based on estimating where are located the set of features that participate the most w.r.t. to the others into supporting a given class. We perform the estimation per class, and the result of this estimation cannot
325 be directly interpreted as a probability. Ultimately, μ is a measure over $\mathbb{F}_{\mathbf{X}}$ directly and it is then obvious it cannot be interpreted as a probability: $\mathbb{F}_{\mathbf{X}} \in 2^{\mathbb{F}_{\mathbf{X}}}$ but by definition of (10), $\mu(\mathbb{F}_{\mathbf{X}}) = 1$ while there is absolutely no reason to think that $J(\mathbb{F}_{\mathbf{X}}) = 1$. Our definition of μ tries to answer the following question: *knowing a certain amount of information materialized by*
330 *the intersection family \mathcal{E} , where are located the features holding more discriminative information than the others?* Once again, with further information about the features, a totally different approach to select μ may be developed and the model (10) offers a more general framework than the particular instance we develop in this paper. To ease the notation and for practical

335 reasons, the measure μ is provided w.r.t. the intersection family \mathcal{E} of the hypergraph induced by the training set \mathbf{X} . However, μ can be defined over any partition of \mathbb{F} .

For \mathbf{x} and \mathbf{x}' in $2^{\mathbb{F}}$, a basic quantitative measure on the importance of \mathbf{x}' w.r.t. \mathbf{x} can be expressed by $\frac{|\mathbf{x} \cap \mathbf{x}'|}{|\mathbf{x}|}$, i.e. how much \mathbf{x}' overlaps with \mathbf{x} . This
 340 measure is a sort of *potential* for an analogy with \mathbf{x} . Potential because it does not account for the individual importance of the features and simply holds the idea that the larger is a subset of features in a case, the higher is the chance it holds important features to decide the outcome.

Let us consider \mathcal{E} and an example $\mathbf{x} \in \mathbf{X}$.

Definition 3.4 (Intrinsic strength w.r.t. \mathbf{x}). The intrinsic strength of $\mathbf{e} \in \mathcal{E}$ w.r.t. $\mathbf{x} \in \mathbf{X}$ is defined by

$$\begin{aligned} \forall l \in \{-1, 1\}, \bar{S}^{(l)}(\mathbf{e}, \mathbf{x}) &= \frac{|d^{(l)}(\mathbf{e})| \frac{|\mathbf{x} \cap \mathbf{e}|}{|\mathbf{x}|}}{\sum_{\mathbf{e}_j \in \mathcal{E}} |d^{(l)}(\mathbf{e}_j)| \frac{|\mathbf{x} \cap \mathbf{e}_j|}{|\mathbf{x}|}} \\ &= \frac{|d^{(l)}(\mathbf{e})| |\mathbf{x} \cap \mathbf{e}|}{\sum_{\mathbf{e}_j \in \mathcal{E}} |d^{(l)}(\mathbf{e}_j)| |\mathbf{x} \cap \mathbf{e}_j|} \end{aligned} \quad (13)$$

Matricially, for any \mathbf{e}_i and \mathbf{x}_j ,

$$\bar{S}_{i,j}^{(l)} = \frac{d_i^{(l)} w_{ji}}{\langle d^{(l)}, w_{i:} \rangle} \quad (14)$$

345 In particular, for a given $\mathbf{x} \in \mathbf{X}$, $\bar{S}^{(l)}(\mathbf{e}_i, \mathbf{x}) = 0$ if \mathbf{e}_i is not part of the projection of \mathbf{x} on H . Also, $\sum_{\mathbf{e} \in \mathcal{E}} \bar{S}^{(l)}(\mathbf{e}, \mathbf{x}) = 1$ which can be interpreted as how much \mathbf{e} accounts for the set of labels l the element \mathbf{x} holds in its projection.

The more \mathbf{e}_i belongs to many cases with the same class l and the higher $\bar{S}^{(l)}(\mathbf{e}_i, \mathbf{x})$ is. Conversely, for a fixed number of cases, the more \mathbf{e}_i describes \mathbf{x} , the higher $\bar{S}^{(l)}(\mathbf{e}_i, \mathbf{x})$ is. As $\forall \mathbf{e}_i \in \mathcal{E}$, $|d(\mathbf{e}_i)| > 0$, either we have $\bar{S}^{(1)}(\mathbf{e}_i, \mathbf{x}) \neq 0$
 350 or $\bar{S}^{(-1)}(\mathbf{e}_i, \mathbf{x}) \neq 0$. We have $\bar{S}^{(l)}(\mathbf{e}_i, \mathbf{x}) = 0$ only when all the cases in which \mathbf{e}_i results of are labeled with the opposite class \bar{l} . For $\bar{S}^{(l)}(\mathbf{e}_i, \mathbf{x}) = 1$, it needs both the unanimity of labels for the cases in which \mathbf{e}_i belongs to and that $\mathbf{e}_i = \mathbf{x}$. The relation $\mathbf{e}_i = \mathbf{x}$ implies that \mathbf{x} does not share any feature with
 355 any other examples or that \mathbf{x} is included into another example.

Definition 3.5 (Intrinsic strength w.r.t. a hypergraph H). The intrinsic strength of $\mathbf{e} \in \mathcal{E}$ w.r.t. $H = (\mathbb{F}_{\mathbf{X}}, \mathbf{X})$ is defined by

$$\begin{aligned} \forall l \in \{-1, 1\}, S^{(l)}(\mathbf{e}) &= \frac{|\mathbf{e}|}{\sum_{\mathbf{e}' \in \mathcal{E}} |\mathbf{e}'|} \sum_{\mathbf{x} \in d^{(l)}(\mathbf{e})} \bar{S}^{(l)}(\mathbf{e}, \mathbf{x}) \\ &= \frac{|\mathbf{e}|}{|\mathbb{F}_{\mathbf{X}}|} \sum_{\mathbf{x} \in d^{(l)}(\mathbf{e})} \bar{S}^{(l)}(\mathbf{e}, \mathbf{x}) \end{aligned} \quad (15)$$

Matricially, for any \mathbf{e}_i ,

$$S_i^{(l)} = \frac{|\mathbf{e}|}{|\mathbb{F}_{\mathbf{X}}|} \|S_{i:}^{(l)}\|_1 \quad (16)$$

The measure $S^{(l)}(\mathbf{e})$ is simply the sum of the relevance for all the examples \mathbf{e} belongs too. The more \mathbf{e} belongs to several cases, the more information it has to support a class or another. As \mathcal{E} represents the sets of features that appear all the time together, we favor the larger $\mathbf{e} \in \mathcal{E}$ as they hold more information to explain a decision. The normalized version is defined by:

$$\forall l \in \{-1, 1\}, \mu^{(l)}(\mathbf{e}) = \frac{S^{(l)}(\mathbf{e})}{\sum_{\mathbf{e}' \in \mathcal{E}} S^{(l)}(\mathbf{e}')} \quad (17)$$

Finally, the measure μ is simply defined by the difference between the strength of both classes:

$$\mu(\mathbf{e}) = \mu^{(1)}(\mathbf{e}) - \mu^{(-1)}(\mathbf{e}) \quad (18)$$

which can be expressed matricially by

$$\mu_i = \frac{|\mathbf{e}|}{|\mathbb{F}_{\mathbf{X}}|} \left[\frac{S_i^{(1)}}{\|S^{(1)}\|_1} - \frac{S_i^{(-1)}}{\|S^{(-1)}\|_1} \right] \quad (19)$$

In a sense, the projection on the hypergraph provides all the possible analogies with the training set, either seen as a support for a particular outcome l or as a counter-example. The intrinsic strength of each element of this partition is a measure of “how much it can be considered as a counter-
360 example or an analogy” taking into account simultaneously all the analogies between the examples that are in relation with the considered case. It favors the importance of the relation (the more features it shares, the stronger is

the analogy) but also the quality of the relation (the more examples sharing the same class the stronger the analogy).

365

Example: Consider the hypergraph in Figure 1 made of \mathbf{x}_1 , \mathbf{x}_2 and \mathbf{x}_3 arbitrarily labeled with resp. 1, -1 and 1. Arbitrarily, we assume the cardinal of the elements of \mathcal{E} to be $\#\mathbf{e} = (2, 1, 2, 3, 1, 2, 3)$ such that the cardinal of cases are $\#\mathbf{x} = (7, 8, 7)$ and $|\mathbb{F}_{\mathbf{x}}| = 14$. The values of $|d^{(l)}(\mathbf{e})|$ can be summarized by the vectors $\mathbf{d}^{(-1)} = (0, 0, 1, 0, 1, 1, 1)$ and $\mathbf{d}^{(1)} = (1, 2, 2, 1, 1, 1, 0)$. Let us calculate $S^{(-1)}(\mathbf{e}_3)$:

370

$$\begin{aligned}\bar{S}^{(1)}(\mathbf{e}_3, \mathbf{x}_1) &= \frac{2 \times 2}{2 \times 1 + 1 \times 2 + 2 \times 2 + 1 \times 2} = \frac{4}{10} \\ \bar{S}^{(1)}(\mathbf{e}_3, \mathbf{x}_2) &= \frac{2 \times 2}{2 \times 2 + 1 \times 1 + 1 \times 2 + 0 \times 3} = \frac{4}{7} \\ \bar{S}^{(1)}(\mathbf{e}_3, \mathbf{x}_3) &= \frac{2 \times 2}{2 \times 1 + 2 \times 2 + 3 \times 1 + 1 \times 1} = \frac{4}{10}\end{aligned}$$

which we interpret as \mathbf{e}_3 being responsible for $\frac{4}{10}$ of the support toward class 1 in \mathbf{x}_1 and \mathbf{x}_3 , while $\frac{4}{7}$ for \mathbf{x}_2 . This lead to

$$S^{(1)}(\mathbf{e}_3) = \frac{2}{14} \left[\frac{4}{10} + \frac{4}{7} + \frac{4}{10} \right] \simeq 0.1959$$

Similarly, we calculate the support for each \mathbf{e} and both labels. We summarize this into the following vectors:

$$\begin{aligned}S^{(1)} &\simeq (0.0286, 0.0286, 0.1959, 0.0643, 0.0173, 0.0694, 0.0000) \\ S^{(-1)} &\simeq (0.0000, 0.0000, 0.2024, 0.0000, 0.0327, 0.1071, 0.0000)\end{aligned}$$

After normalization, we obtain the intrinsic strength:

$$\mu \simeq (0.0707, 0.0707, 0.0060, 0.1591, -0.0345, -0.0818, -0.1901)^T$$

Let us evaluate the model on the three examples:

$$\begin{aligned}s(\mathbf{x}_1) &= \frac{2}{7}\mu(\mathbf{e}_1) + \frac{1}{7}\mu(\mathbf{e}_2) + \frac{2}{7}\mu(\mathbf{e}_3) + \frac{2}{7}\mu(\mathbf{e}_6) \simeq 0.0086 \\ s(\mathbf{x}_2) &= \frac{2}{8}\mu(\mathbf{e}_3) + \frac{1}{8}\mu(\mathbf{e}_5) + \frac{2}{8}\mu(\mathbf{e}_6) + \frac{3}{8}\mu(\mathbf{e}_7) \simeq -0.0946 \\ s(\mathbf{x}_3) &= \frac{1}{7}\mu(\mathbf{e}_2) + \frac{2}{7}\mu(\mathbf{e}_3) + \frac{3}{7}\mu(\mathbf{e}_4) + \frac{1}{7}\mu(\mathbf{e}_5) \simeq 0.0751\end{aligned}$$

As a result, \mathbf{x}_1 and \mathbf{x}_3 are labeled 1 and \mathbf{x}_2 is labeled -1 . All three cases are
 375 correctly labeled. The highest support is given for case \mathbf{x}_2 and \mathbf{x}_3 while the
 support for \mathbf{x}_1 is one order of magnitude lower then for \mathbf{x}_3 . This is because
 the discretionary features of \mathbf{x}_3 are larger while the intersection with \mathbf{x}_2 is
 lower than for \mathbf{x}_1 ($\frac{3}{8}$ of \mathbf{x}_3 against $\frac{4}{7}$ of \mathbf{x}_1).

Consider a new case \mathbf{x} as described in Figure 2. Its support is given by
 380 $s(\mathbf{x}) = \sum_{\mathbf{e} \in \pi(\mathbf{x})} w(\mathbf{e}, \mathbf{x}) \mu(\mathbf{e})$ with $\pi(\mathbf{x}) = \{\mathbf{e}_1, \mathbf{e}_2, \mathbf{e}_3, \mathbf{e}_5, \mathbf{e}_6, \mathbf{e}_7\}$. It is impossible
 for \mathbf{x} to be classified as 1 because the highest support would be for a maximal
 intersection with $\mathbf{e}_1, \mathbf{e}_2, \mathbf{e}_3$ and minimal for $\mathbf{e}_5, \mathbf{e}_6$ and \mathbf{e}_7 such that $s(\mathbf{x}) =$
 $\frac{1}{8}(2\mu(\mathbf{e}_1) + \mu(\mathbf{e}_2) + 2\mu(\mathbf{e}_3) + \mu(\mathbf{e}_5) + \mu(\mathbf{e}_6) + \mu(\mathbf{e}_7)) \simeq -0.0103 < 0$. It can
 be explained by the fact that the support for 1 is provided by a larger set
 385 of features (11 features versus 8). On top of that, the intersections between
 positive cases (\mathbf{e}_2 and \mathbf{e}_3) are too small (1 for \mathbf{e}_2 compared to e.g. 3 for \mathbf{e}_7)
 or include also negative cases (\mathbf{e}_3).

3.4. Training and Decision

At this stage, it is already possible to generate predictions for new cases.
 390 However, the intrinsic strength vector calculated on the hypergraph might not
 be perfectly accurate on the training set because of the lack of information
 contained in the training set (or some limitations on the model space itself
 that we will discuss in Section 4.2). In this section, we give an algorithm to
 adjust the intrinsic strength in order to correct its initial estimation and we
 395 update the classification rule to take into account the fact the model cannot
 provide predictions over \mathbb{F} entirely.

Once the model is built, it can be evaluated on the training set. Anal-
 ogously to SVM, we define the margin as the distance to the correct class,
 i.e. $m(w, \mu, \mathbf{x}) = J(\mathbf{x})_{s_{w, \mu}(x)}$. To improve the pertinence of the strength of
 400 the elements of \mathcal{E} , we use the iterative algorithm described by Algorithm 2
 to minimize total margin over the training set \mathbf{X} . When and only when a
 classification for a case \mathbf{x} is wrong, it modifies each element of its projection
 by lowering its strength for the wrong class and increasing it for the proper
 class. The margin is split between the element of the projection w.r.t. their
 405 respective weight in \mathbf{x} i.e. $w(\mathbf{e}, \mathbf{x})$. If a case \mathbf{x} is wrongly classified, it is due to
 the cases intersecting with it. Indeed, if \mathbf{x} was not intersecting with any other
 example, its projection would be itself, and its support toward the wrong
 class would be 0 and positive for the real class. In other words, \mathbf{x} would
 be correctly classified. Thus, the idea is not to directly bring the support

410 of x to the correct class but to gradually adjust the weights such that the neighbors are modified enough for \mathbf{x} to be correctly classified. In particular, it is sensitive to the order in which the cases are considered: a modification in the strength of any $\mathbf{e} \in \mathcal{E}$ impacts all cases in which it appears and potentially changes the predicted class for those cases. In addition, there is no guarantee
 415 of convergence. Future work will focus on the optimal order or a modification schema such that the algorithm converges. Investigating other minimizing functions is also considered. However, empirical experiments (see Section 4.2) suggests that the result after training is near-optimal w.r.t. to performances observed on the test set.

Algorithm 2 Model training

Input:

- \mathbf{X} : training set
- \mathbf{y} : correct labels for \mathbf{X}
- k : number of training iterations
- $\mu^{(1)}, \mu^{(-1)}$: weights calculated with (3.5)

Output:

- Modified vectors $\mu^{(1)}, \mu^{(-1)}$

```

1: for  $k$  iterations do
2:   for  $\mathbf{x}_i \in \mathbf{X}$  do
3:      $\bar{y}_i \leftarrow \bar{J}(\mathbf{x}_i)$ 
4:     if  $\bar{y}_i \neq y_i$  then
5:       for  $\mathbf{e} \in \pi(\mathbf{x}_i)$  do
6:          $\mu^{(y_i)}(\mathbf{e}) \leftarrow \mu^{(y_i)}(\mathbf{x}_i) + w(\mathbf{e}, \mathbf{x}_i)|\mu(\mathbf{e})|$ 
7:          $\mu^{(\bar{y}_i)}(\mathbf{e}) \leftarrow \mu^{(\bar{y}_i)}(\mathbf{x}_i) - w(\mathbf{e}, \mathbf{x}_i)|\mu(\mathbf{e})|$ 
8:       end for
9:     end if
10:   end for
11: end for

```

The measure μ is defined as the difference of support for both classes. Thus, by linearity we can rewrite

$$\begin{aligned}
 s(\mathbf{x}) &= \sum_{i=1}^M w(\mathbf{e}_i, \mathbf{x}) \mu^{(1)}(\mathbf{e}_i) - \sum_{i=1}^M w(\mathbf{e}_i, \mathbf{x}) \mu^{(-1)}(\mathbf{e}_i) \\
 &= s^{(1)}(\mathbf{x}) - s^{(-1)}(\mathbf{x})
 \end{aligned} \tag{20}$$

This form is convenient because we can control how much evidence we need to support a specific class using the following constraints and a family η of four hyperparameters:

$$s^{(-1)}(\mathbf{x}) > \max\left(\frac{\bar{\eta}_{-1}}{1 - \bar{\eta}_{-1}} s^{(-1)}(\mathbf{x}), \eta_{-1}\right) \geq 0 \quad (C_0)$$

$$s^{(1)}(\mathbf{x}) > \max\left(\frac{\bar{\eta}_1}{1 - \bar{\eta}_1} s^{(1)}(\mathbf{x}), \eta_1\right) \geq 0 \quad (C_1)$$

420 with $\eta_{-1}, \eta_1 \in \mathbb{R}^+$ and $\bar{\eta}_0, \bar{\eta}_1 \in [0, 1]$. The constraints on η_{-1} and η_1 defines a minimal amount of support respectively toward class -1 and 1 while $\bar{\eta}_{-1}$ and $\bar{\eta}_1$ requires the support toward a class to be significantly higher than the support for the other class. As μ is normalized over the partition \mathcal{E} , the value of η_{-1} and η_1 must be set w.r.t. the hypergraph. On the contrary, $\bar{\eta}_1$ and $\bar{\eta}_0$
425 can be set independently of the hypergraph.

Those constraints may be used to design a decision rule for new cases depending on the application or the dataset. The most generic decision rule is as follows:

$$\tilde{J}(\mathbf{x}) = \begin{cases} 1 & s(\mathbf{x}) > 0 & \text{and } C_1 \\ 0 & s(\mathbf{x}) \leq 0 & \text{and } C_0 \\ l_1 & s(\mathbf{x}) > 0 & \text{and not } C_1 \\ l_{-1} & s(\mathbf{x}) \leq 0 & \text{and not } C_0 \end{cases} \quad (22)$$

where l_{-1}, l_0 are two labels. A representation is given by Figure 3. Those hyperparameters are intended to model the “burden of proof”. For instance, in a trial, one is assumed innocent until proven guilty which implies the support for the class “guilty” must be *beyond a reasonable doubt* (where the
430 term reasonable is defined by the jurisprudence of the applicable country). In case $\eta_{-1} = \eta_1 = \bar{\eta}_{-1} = \bar{\eta}_1$ (and $l_{-1} = 0$ and $l_1 = 1$), then the decision rule is equivalent to the original one defined by (R1). In case \mathbf{x} has too many discretionary features, this classification rule is likely to be irrelevant. Indeed, the intersection between \mathbf{x} and $\mathbb{F}_{\mathbf{x}}$ is too small to hold enough information
435 and make strong analogies with \mathbf{x} . To overcome this drawback, $2^{\mathbb{F}}$ is split into two subsets:

- $\mathcal{F}_1 = \{\mathbf{x} \in 2^{\mathbb{F}} \mid |\mathbf{x} \cap \mathbb{F}_{\mathbf{x}}| \geq \delta\}, \forall \delta \in \mathbb{N}$
- $\mathcal{F}_2 = 2^{\mathbb{F}} \setminus \mathcal{F}_1$




Figure 3: Representation of the updated decision rule (22).

\mathcal{F}_1 corresponds to the elements such that they share some features with the examples. An alternative may be considered by using $\mathcal{F}_1 = \{\mathbf{x} \in 2^{\mathbb{F}} \mid \frac{D_{\mathbf{x}}}{|\mathbf{x}|} \leq \delta\}$, $\forall \delta \in [0, 1]$. In this case, \mathcal{F}_1 contains the elements for which we have enough information provided by the examples. From our preliminary tests, the choice depends on the dataset structure.

Finally, the decision rule for new cases is built as follows:

$$\bar{J}(\mathbf{x}) = \begin{cases} \tilde{J}(\mathbf{x}) & \text{if } \mathbf{x} \in \mathcal{F}_1 \\ o_{\mathbf{x}} & \text{if } \mathbf{x} \in \mathcal{F}_2 \end{cases} \quad (\text{R2})$$

where $o_{\mathbf{x}}$ is one draw from a random variable that has Bernoulli law with parameter $p = \frac{|\{\mathbf{x} \in \mathbf{X} \mid J(\mathbf{x})=1\}|}{|\mathbf{X}|}$, i.e. the prevalence of 1 in \mathbf{X} . This assumes that \mathbf{X} is correctly representing $2^{\mathbb{F}}$ (or that the prevalence does not change in time for sequential problems in which the new cases are generated by an unknown random measure). The rational behind is that if for a case \mathbf{x} , it is not possible to exploit the model built on the hypergraph, then we can still consider that J acts as a Bernoulli random variable and use a maximum likelihood estimation for p . In a sense, it is extending the *local* model to the entire input space $2^{\mathbb{F}}$.

3.5. Complexity

Model Building: Given $\mathbf{X} \in (2^{\mathbb{F}})^N$, constructing \mathcal{E}_H can be done in $\mathcal{O}(\sum_{\mathbf{x} \in \mathbf{X}} |\mathbf{x}|)$ by using a Partition Refinement data structure [6]. Given $\mathbf{x} \in \mathbf{X}$, calculating the family $\{\bar{S}(\mathbf{e}, \mathbf{x})\}_{\mathbf{e} \in \mathcal{E}_H}$ can be done in $\mathcal{O}(|\mathbf{x}|)$ by asking for each feature of \mathbf{x} the \mathbf{e} it belongs to and maintaining the size of each \mathbf{e} during

the construction of \mathcal{E}_H . Thus, calculating $\{\bar{S}(\mathbf{e}, \mathbf{x})\}_{\mathbf{e} \in \mathcal{E}_H}$ for all $\mathbf{x} \in \mathbf{X}$ can be done in $\mathcal{O}(\sum_{\mathbf{x} \in \mathbf{X}} |\mathbf{x}|)$. On m -uniform hypergraphs (when all cases are described
460 with m features) with n hyperedges, it becomes $\mathcal{O}(mM)$.

Calculating $\{S(\mathbf{e})\}_{\mathbf{e} \in \mathcal{E}_H}$ and μ can be done in $\mathcal{O}(|\mathcal{E}_H|)$ because it requires to iterate over \mathcal{E}_H . An obvious upper bound on $|\mathcal{E}_H|$ is $|\mathbb{F}_{\mathbf{X}}|$ i.e. the number of vertices in the hypergraph. The worst-case cardinal of \mathcal{E}_H is when each $\mathbf{x} \in \mathbf{X}$ intersects with all the others and none of them is strictly a subset of
465 any other. Thus, $|\mathcal{E}_H| \leq \min(2^N - 1, |\mathbb{F}_{\mathbf{X}}|)$.

Learning Phase: For each wrongly classified $\mathbf{x} \in \mathbf{X}$, the training requires at most $\mathcal{O}(|\mathbf{x}|)$ steps (maximal cardinal for $\pi(\mathbf{x})$). The worst-case scenario is when the model wrongly classifies every $\mathbf{x} \in \mathbf{X}$. Thus, the learning phase
470 worst-case complexity is $\mathcal{O}(k \sum_{\mathbf{x} \in \mathbf{X}} |\mathbf{x}|)$ and on m -uniform hypergraphs it becomes $\mathcal{O}(kmN)$.

Model Query: For a case $\mathbf{x} \in 2^{\mathbb{F}}$, the projection can be done in $\mathcal{O}(|\mathbf{x}|)$. Calculating the classification rule also requires at most $\mathcal{O}(|\mathbf{x}|)$ (maximal
475 cardinal for $\pi(\mathbf{x})$).

4. Experiments

The goal of the experiments is twofold: validating the practical interest of HCBR w.r.t. the literature results and other classification methods, and studying important properties for practical situations such as the computation
480 time or influence of parameters. The experiments are broken down into two parts: the comparison with literature results and standard alternative methods in terms of classification performance in Section 4.1, and intrinsic performance (computation time, learning curves, model space limitations and influence of parameters) in Section 4.2.

For the first part, we reported the confusion matrix obtained over all the runs but also after each prediction. From this confusion matrix, we calculated the standard performance indicators: accuracy, recall, specificity, precision, negative prediction value, F_1 -score and Matthews correlation coefficient. Denoting by TP the number of true positives, TN the true negatives, FP the false positives and FN the false negative, the F_1 -score and Matthews

correlation coefficient are defined by:

$$F_1 = \frac{2TP}{2TP + FP + FN}$$

$$MCC = \frac{TP \times TN - FP \times FN}{\sqrt{(TP + FP)(TP + FN)(TN + FP)(TN + FN)}}$$

F_1 score and Matthews correlation coefficients respectively belongs to $[0, 1]$ and $[-1, 1]$ and the closer to 1, the better it is. Both takes into account false positive and false negatives. In particular, Matthews correlation coefficient is very adapted for binary classification on unbalanced dataset (with a prevalence far from 0.5). However, as many studies do not report them, we will base our comparison with literature results on the accuracy defined by:

$$ACC = \frac{TP + TN}{TP + TN + FP + FN}$$

485 For the second part, we studied the computation time depending on the number of cases and the size of each case in order to validate the worst-case complexity given in Section 3.5. We also studied the learning curve to determine if the model tends to have large variance or bias and the intrinsic limits of the model space limitation, as well as the influence of
490 hyperparameters.

The integrality of the data used for the experiments, as well as the scripts to transform them and analyze the results are available within the HCBR Github repository⁵ and the whole experimental campaign starting from the raw data can be reproduced in "one click".

495 4.1. Classification performances

4.1.1. Literature comparison

To validate the approach, we measured the accuracy, F1-score and Matthew Correlation Coefficient for 7 datasets for binary classification have been used. They are available either from the UCI Machine Learning Repository⁶ or
500 provided with the LIBSVM⁷ : **adult**, **breasts**, **heart**, **mushrooms**, **phishing**,

⁵<https://github.com/aquemy/HCBR>

⁶<https://archive.ics.uci.edu/ml/index.php>

⁷<https://www.csie.ntu.edu.tw/~cjlin/libsvmtools/datasets/binary.html>

Table 1: Datasets description.

	Cases	Total Features	Unique	Min. Size	Max. Size	Average Size	Real
adult	32561	418913	118	10	13	12.87	No
breasts	699	5512	80	8	8	8	No
heart	270	3165	344	12	13	12.99	Yes
mushrooms	8124	162374	106	20	20	20	No
phishing	11055	319787	808	29	29	29	No
skin	245057	734403	768	3	3	3	Yes
splice	3175	190263	237	60	60	60	No

skin and **splice**. For each dataset, the original features (**name=value**) are converted into a unique identifier and the union of all such identifiers constitute the information set \mathbb{F} considered by the algorithm. Notice that there is no need to remove the rows with empty values. Table 1 describes each dataset. The minimal, maximal and average size give information about the case sizes (notice some cases are missing values for **adult**, **heart** and **mushrooms** datasets). The unique features are the number of (**name=value**) in the original dataset. In addition, two datasets have at least one real-valued attribute as indicated by the column "Real" in Table 1.

The validation was made using a 10-fold cross-validation: each dataset has been split into 10 equal sized samples, 90% has been used as training set and the remaining 10% to test the classification. Each subsample is used once as testing set and the final metrics are calculated as the average of the 10 runs. The training set was not rebalanced and kept the original prevalence. The number of training steps k was adjusted with a manual trial and errors approach, the family η set to $\eta_{-1} = \eta_1 = \bar{\eta}_{-1} = \bar{\eta}_1 = 0$ and δ was set to 1. Further work will focus on automatic parameter tuning. The average confusion matrix obtained for each dataset is showed in Table 2. The performance indicators are reported in Table 3. The proposed algorithm performs very well on a wide range of datasets as reported by tables 2 and 3, in particular when they contain strong predictors as it is the case for **mushroom**. The accuracy is contained in a range from 0.8206 (**adult**) to 1 (**mushrooms**) while the F_1 -score is bounded by 0.8653 (**heart**) and 1 (**mushrooms**). On **adult**, the accuracy is only 6% higher than the prevalence, i.e. a baseline model consisting in returning 1 for any point would be only 6% worse. This

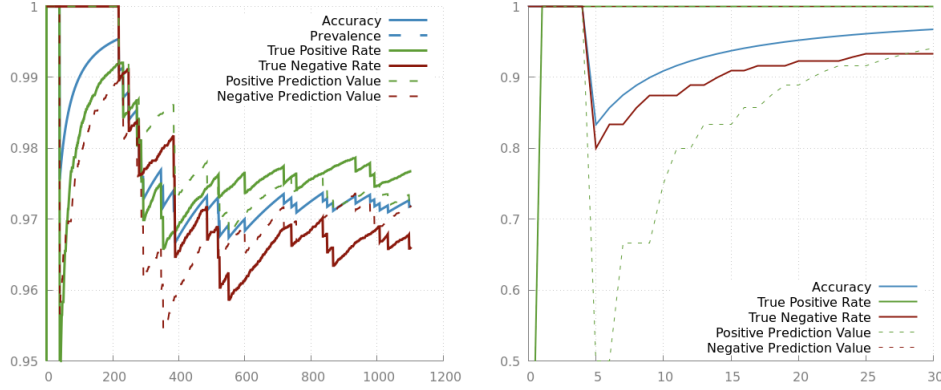


Figure 4: Evolution of the confusion matrix during the prediction phase for **phishing** and **splice** datasets.

relatively poor performance in learning the underlying decision mapping is better reflected by the Matthews correlation coefficient of 0.51.

By default, the false positives and false negatives are equilibrated for each dataset, despite a huge variation in the prevalence (between 20% and 64%, cf. Table 2) which may be a desirable property depending on applications. In addition, the results were obtained with non-balanced training sets which are known to be a problem for many machine learning algorithms [7]. Moreover, the performance indicators remain stable during the whole prediction phase as shown by Figure 4 for the dataset **phishing** (similar result is observed for all datasets).

For a given case, the support is a metric of confidence in the prediction as illustrated in Figure 5. In general, the wrongly classified cases have a smaller difference between the evidence for each class which confirm the interest in the hyperparameters η and $\bar{\eta}$ used in (R2). This can also be observed in Figure 6.

To compare the results of the proposed method, we explored the best results from the literature for each dataset. The results are summarized in Table 5. In [8], 5 rule-based classification techniques dedicated to medical databases are compared and achieve at best 95.85% and 82.96% accuracy resp. on **breast**, and **heart** datasets. Comparing bayesian approaches, [9] demonstrated 97.35% (**breast**) and 83.00% (**heart**) accuracy. A 5 layers neural network with fuzzy inference rules achieved 87.78% on **heart** [10] while a k-NN algorithm reached 99.96% on **mushrooms** [11]. The best alternative

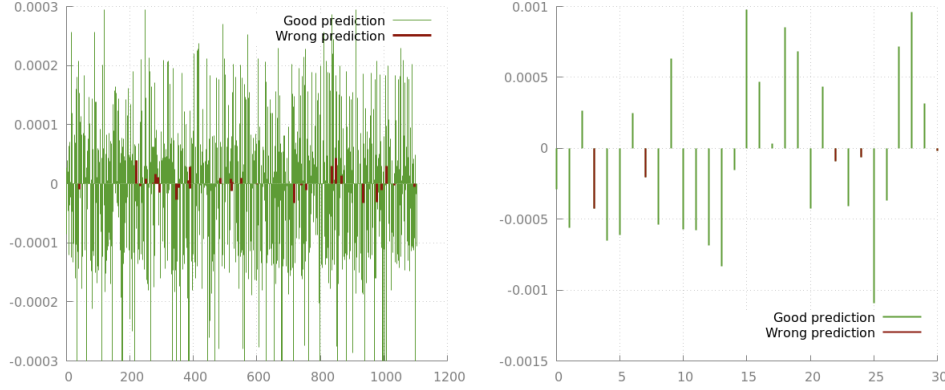


Figure 5: Difference between the weight assigned to both classes for each decision on **phishing** and **splice** (average). Similar results are observed for all datasets.




Figure 6: Histogram of decisions depending on the strength for **phishing** and **splice**. In blue the correctly classified elements, in red the wrongly classified ones. The false positives and false negatives are concentrated around 0. Similar results are observed for all datasets.

among 6 rules-based classification methods achieved 95.84% on **breast** and
 550 100.00% on **mushroom** [12]. Using 80% of **phishing** as training set, an
 adaptive neural network achieved an average accuracy of 93.76% (among
 6 alternatives) with the best run at 94.90% [13]. Still on **phishing**, [14]
 proposes to combine several classifiers and reaches 97.75% accuracy for the
 best hybrid model (and demonstrates 97.58% for Random Forest classifier).
 555 On **adult**, the comparison of several classifiers (naive bayes, decision tree, ...)
 demonstrated at most 86.25% accuracy [15] while a Support Vector Machine
 approach reached 85.35% [16]. On **splice**, a method using Fuzzy Decision
 Trees [17] reaches 94.10% accuracy and a neural network combined to boosting

[18] 97.54%. On **breast**, Support Vector Machine approaches reached resp.
560 96.87%, 98.53%, 99.51% accuracy [19, 20, 21], 99.26% and 97.36% for neural
network based techniques [22, 23], 98.1% for a bayesian network method [24],
or 94.74% using Decision Trees [25]. On **skin**, [18] reports 98.94% accuracy
against 99.68% for Decision Tree based method [26]. The best result, as far
as we know, is 99.92%, obtained by a Generalized Linear Model [27] (with
585 80% training set).

The accuracy is slightly lower than the best results from the literature
(**adult** 82.06% against 86.25%, **breast** 96.96% against 99.51%, **heart** 85.77%
against 87.78%, **phishing** 96.05% against 97.75%, **splice** 94.43% against
97.54%, **skin** 98.65% against 99.92%). We explain this by at least two factors.
570 First, the best methods on a given dataset are often dedicated to this dataset
with *ad-hoc* or engineered parts which is not the case of HCBR. Secondly, the
hyperparameter family η have not been tuned but as we will show in next
section, they might have a decisive impact on performance. HCBR performed
better than Bayes classifier in two thirds of cases. Bayes classifier performs
575 better on **breast** by approximately 1% which represents less than one case
wrongly classified. Similar results are observed with Decision Trees. However,
the 1% difference on **skin** represents an average of 7 cases misclassified in
comparison in favor of Bayes. It performs better than Rule-based approaches
(or gives similar results on **mushrooms** with an accuracy of 1) in the four
580 considered references on three different datasets. Notice that combined with
Neural Network, Rule-based achieves the state-of-art result on **heart** dataset.
Except for **phishing**, Neural Network returns better results (0.46 more cases
with correct classification in average for **breast**, 71 for **skin** and almost 10
for **splice**). Last, SVM gives better results in all three cases, but appear
585 only as best results in two datasets.

Even if HCBR does not rank first, we see at least three reasons to use
it in practice. First, it provides consistently good results on all datasets
without a need to tune the hyperparameters (which will be confirmed with
the second round of experiments below), without using stratified sample to
590 correct the unbalanced datasets or feature transformation that would require
expert knowledge. In other words, it requires absolutely no domain knowledge
or data science experience to deploy and use. Second, HCBR works directly
on unstructured dataset without a need of some exotic representations or
encoders as it is the case for other methods. This has a considerable advantage
595 in practice where the information per case is rarely structured by default.
Last but not least, HCBR provides local explanation: for each case and each

group of features \mathbf{e}_i in this case, HCBR provides not only the support for \mathbf{e}_i but also all the cases that participated in increasing the support. In the medical domain, the practitioner may check for instance the top two groups of features and for each the five most influencing past cases to find a reasonable justification to the prediction.

Table 2: Average confusion matrix obtained with a 10-fold cross-validation.

	TP	FN	FP	TP	Prevalence
adult	2182.40	295.30	288.50	488.80	0.7586
breast	23.00	1.40	0.70	43.90	0.3338
heart	12.40	1.80	1.90	9.90	0.5107
mushrooms	390.60	0.00	0.00	420.40	0.4804
phishing	595.40	23.80	19.80	465.00	0.5562
skin	4886.30	132.40	199.40	19286.90	0.2075
splice	155.70	9.10	8.50	142.70	0.5164

Table 3: Average performances obtained with a 10-fold cross-validation.

	Accuracy (std dev.)	Recall	Specificity	Precision	Neg. Pred. Value	F ₁ score	MCC
adult	0.8206 (0.0094)	0.8832	0.6233	0.8808	0.6290	0.8820	0.5081
breasts	0.9696 (0.0345)	0.9691	0.9676	0.9479	0.9844	0.9575	0.9344
heart	0.8577 (0.0943)	0.8695	0.8437	0.8699	0.8531	0.8653	0.7178
mushrooms	1.0000 (0.0000)	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000
phishing	0.9605 (0.0081)	0.9680	0.9514	0.9615	0.9590	0.9647	0.9199
skin	0.9865 (0.0069)	0.9608	0.9932	0.9736	0.9898	0.9672	0.9587
splice	0.9443 (0.0124)	0.9478	0.9398	0.9450	0.9441	0.9463	0.8884

4.1.2. Robustness comparison

As mentionned above, the state-of-art results obtained per database required a lot of efforts in terms of feature selection and transformation, as well as model selection or hyperparameter tuning. Therefore, the obtained models are slightly better than HCBR but not well adapted to solve several instances of the classification problem. If well-calibrated models with high performances on their domain of application is undeniably useful for practical usage, being able to solve a problem on a large variety of instances without

Table 4: Previous literature results measured as the highest accuracy obtained by the authors.

Dataset	Ref.	Type	Accuracy
adult	[15]	Many classifiers	86.25%
	[16]	SVM	85.35%
		HCBR	82.06%
breast	[21]	SVM	99.51%
	[22]	Neural Net	99.26%
	[20]	SVM	98.53%
	[24]	Bayes	98.1%
	[23]	Neural Net	97.36%
	[9]	Bayes	97.35%
		HCBR	96.96%
	[19]	SVM	96.87%
	[8]	Rule-based	95.85%
	[12]	Rule-based	95.84%
	[25]	Decision Tree	94.74%
heart	[10]	Neural Network + Rule-based	87.78%
		HCBR	85.77%
	[9]	Bayes	83.00%
	[8]	Rule-based	82.96%
mushrooms	[12]	Rule-Based	100.00%
		HCBR	100.00%
	[11]	k-NN	99.96%
phishing	[14]	Ensemble	97.75%
	[14]	Random-Forest	97.58%
		HCBR	96.05%
	[13]	Neural Net	94.90%
skin	[27]	Generalized Linear Model	99.92%
	[26]	Decision Tree	99.68%
	[18]	Neural Network + Boosting	98.94%
		HCBR	98.65%
splice	[18]	Neural Network + Boosting	97.54%
		HCBR	94.43%
	[17]	(fuzzy) Decision Tree	94.10%

610 effort is also of a vital importance, especially in the industry where the end-user might not be specialized in data science and machine learning. One main

challenge of large scale machine learning systems is thus to achieve this good compromise between ready-to-deploy models and good performance metrics. With this perspective, comparing the state-of-art result per instance might
615 not be a good indicator of this scalability capacity of a particular method.

For a fair comparison, we performed a 10 cross-validation for each of the 7 selected datasets using 9 standard classification methods: AdaBoost, k-Nearest Neighbors, Linear SVM, Radius-Based Function (RBF) SVM, Decision Tree, Random Forest, Neural Network and Quadratic Discriminant
620 Analysis (QDA). The implementation is provided by Scikit-Learn [28]. No hyperparameter tuning was performed and the default values of parameters were used. Matthew Correlation Coefficient has been shown to be more informative than other metrics derivated from the confusion matrix [29], in particular with unbalanced datasets. For this reason, we used it instead of
625 the accuracy or F_1 score and reported the MCC per dataset as well as the average MCC and the corresponding rank (see Table 6). Additionally, we displayed the evolution of MCC with the training set size in Figure 7.

The highest MCC value is obtained by Neural Network with 0.8914 followed by HCBR with 0.8435. Neural Network improves HCBR result by 5.68%
630 while HCBR improves the result of all other methods from 2% to 17.56%. The lowest MCC score is obtained by Naive Bayes with 0.6953, mostly because it performs poorly on some datasets (0.2493 on `adult`, 0.5292 on `phishing` or 0.7600 on `skin`). In general, Neural Network and HCBR are the only two methods whose rank remains consistent across all datasets, close to the
635 first and second rank. On the contrary, methods like k -Nearest Neighbors or Decision Tree performed very well on one dataset (`skin`, resp. `splice`). Surprisingly, Random Forest not only perform poorer than Decision Tree but also performs worse than most methods. Naive Bayes and QDA perform very poorly in general which is less surprising knowing the assumptions behind
640 those methods that are likely to be unrealistic on real datasets.

In other words, HCBR and Neural Network have shown to be more versatile than the other approaches on this selected set of classification instances and without parameter tuning. We have no doubt that the results obtained with other methods can be highly improved by a proper hyperparameter
645 selection, in particular for Random Forest or even k-Nearest Neighbors.

4.2. Intrinsic performance

In this section, we validate the time complexity and discuss the model space limitations of HCBR. We also study the influence of the two main

Table 5: Comparison of HCBR with several methods (Scikit-Learn implementation).

Dataset	Method	MCC	#
adult	HCBR	0.5146	3
	AdaBoost	0.5455	1
	k-NN	0.4785	7
	Linear SVM	0.4918	5
	RBF SVM	0.5065	4
	Decision Tree	0.4821	6
	Rand. Forest	0.3776	8
	Neural Net	0.5349	2
	Naive Bayes	0.2493	9
	QDA	0.4785	7
breast	HCBR	0.9222	3
	AdaBoost	0.9023	6
	k-NN	0.9163	4
	Linear SVM	0.9126	5
	RBF SVM	0.8829	8
	Decision Tree	0.8760	9
	Rand. Forest	0.9296	1
	Neural Net	0.9280	2
	Naive Bayes	0.8991	7
	QDA	0.8616	10
heart	HCBR	0.7082	1
	AdaBoost	0.5972	6
	k-NN	0.5879	7
	Linear SVM	0.6849	4
	RBF SVM	0.6287	5
	Decision Tree	0.5763	8
	Rand. Forest	0.5703	9
	Neural Net	0.6995	2
	Naive Bayes	0.6932	3
	QDA	0.4500	10
mushrooms	HCBR	0.9995	2
	AdaBoost	1.0000	1
	k-NN	0.9993	3
	Linear SVM	1.0000	1
	RBF SVM	0.9990	5
	Decision Tree	0.9991	4
	Rand. Forest	0.8840	7
	Neural Net	1.0000	1
	Naive Bayes	0.9767	6
	QDA	1.0000	1

Dataset	Method	MCC	#
phishing	HCBR	0.9191	1
	AdaBoost	0.8637	6
	k-NN	0.9138	4
	Linear SVM	0.8740	5
	RBF SVM	0.9286	2
	Decision Tree	0.8585	7
	Rand. Forest	0.7582	8
	Neural Net	0.9448	1
	Naive Bayes	0.5292	10
	QDA	0.5872	9
skin	HCBR	0.9551	4
	AdaBoost	0.8552	8
	k-NN	0.9982	1
	Linear SVM	0.8090	9
	RBF SVM	0.9950	3
	Decision Tree	0.9544	5
	Rand. Forest	0.9539	6
	Neural Net	0.9967	2
	Naive Bayes	0.7600	10
	QDA	0.9483	7
splice	HCBR	0.8857	2
	AdaBoost	0.8801	3
	k-NN	0.6072	9
	Linear SVM	0.7282	8
	RBF SVM	0.8461	4
	Decision Tree	0.8998	1
	Rand. Forest	0.5925	10
	Neural Net	0.8390	5
	Naive Bayes	0.7595	7
	QDA	0.8251	6

Table 6: Average MCC and rank obtained with several methods (Scikit-Learn implementation). The column Δ_{HCBR} represents the relative difference in MCC w.r.t. to HCBR.

Method	MCC	Rank	Δ_{HCBR}
Neural Network	0.8914	1	5.68%
HCBR	0.8435	2	-
RBF SVM	0.8267	3	2.00%
Decision Tree	0.8066	4	4.37%
AdaBoost	0.8063	5	4.41%
k-NN	0.7859	6	6.82%
Linear SVM	0.7858	7	6.84%
QDA	0.7358	8	12.76%
Random Forest	0.7237	9	14.20%
Naive Bayes	0.6953	10	17.56%

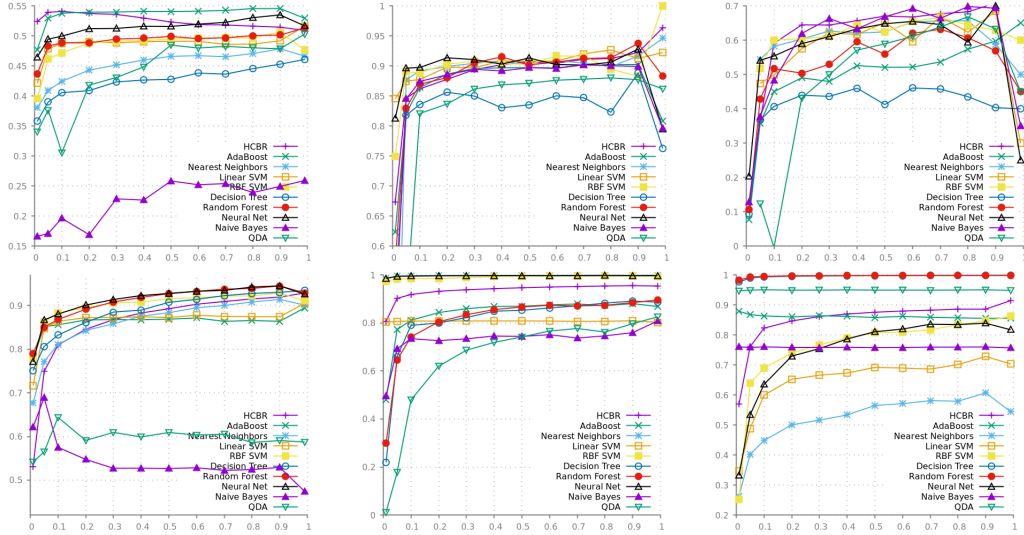


Figure 7: Matthew Correlation Coefficient depending on the training set size for **adult**, **breast**, **heart**, **phishing**, **skin** and **splice**. The fitness is lower than the MCC by definition of (28).

parameters (number of examples and size of the examples) on the computation time.

4.2.1. Computation Time

We generated a casebase of N cases of size m such that case i is described by $\{i, \dots, i+m\}$ i.e., each case is partitioned into m elements (one discretionary feature). This is the worst-case scenario in terms of the size of \mathcal{E} if $m < N$ because the family grows exponentially in function of m or N . We split the computation time into constructing the hypergraph (and determining the intersection family) and calculating the strength of the partition. The results are illustrated in Figure 8. By increasing N with a fixed m , the partition grows exponentially and thus, it is expected to have an exponential curve for the strength computation. On the contrary, building the hypergraph can be done in linear time when m fixed. When N is fixed and m increases, constructing the hypergraph is still doable in linear time as expected. Interestingly, calculating the strength has two phases: if $m \leq N$, increasing m exponentially increases the time (because \mathcal{E} exponentially increases) but if $m > N$, increasing m cannot results in an exponential growth in the computation time (because \mathcal{E} grows linearly).




Figure 8: At the top, computation time to build the model (hypergraph construction + strength calculation) depending on N ($m = 10$), and at the bottom, depending on m ($N = 100$). The case i is described by $\{i, \dots, i+m\}$ such that each case is partitioned into m elements (one discretionary feature). Right scale for building and left scale for strength.

4.2.2. Learning Curves

The learning curves are useful to study the limit of the model space on the different datasets. It consists in plotting the accuracy in function of the training set size for both the predictions over the training and the test sets.

For the training set, it is expected to observe an accuracy starting close to 1 and decreasing fast to reach a plateau. A low stationary accuracy value indicates a high bias in the model or/and an irreducible error contained in the dataset, such as noisy or uninformative features. Oppositely, the accuracy on the test sets starts close to 0 as the training set is small and should increase until a plateau which is very often expected to be lower than of the training set. If the training set curve converges toward a much higher value than the test set curve, then the model has a large variance. In other words, to achieve a good bias-variance tradeoff, the accuracy of both curves should converge toward more or less the same value, expected as high as possible.

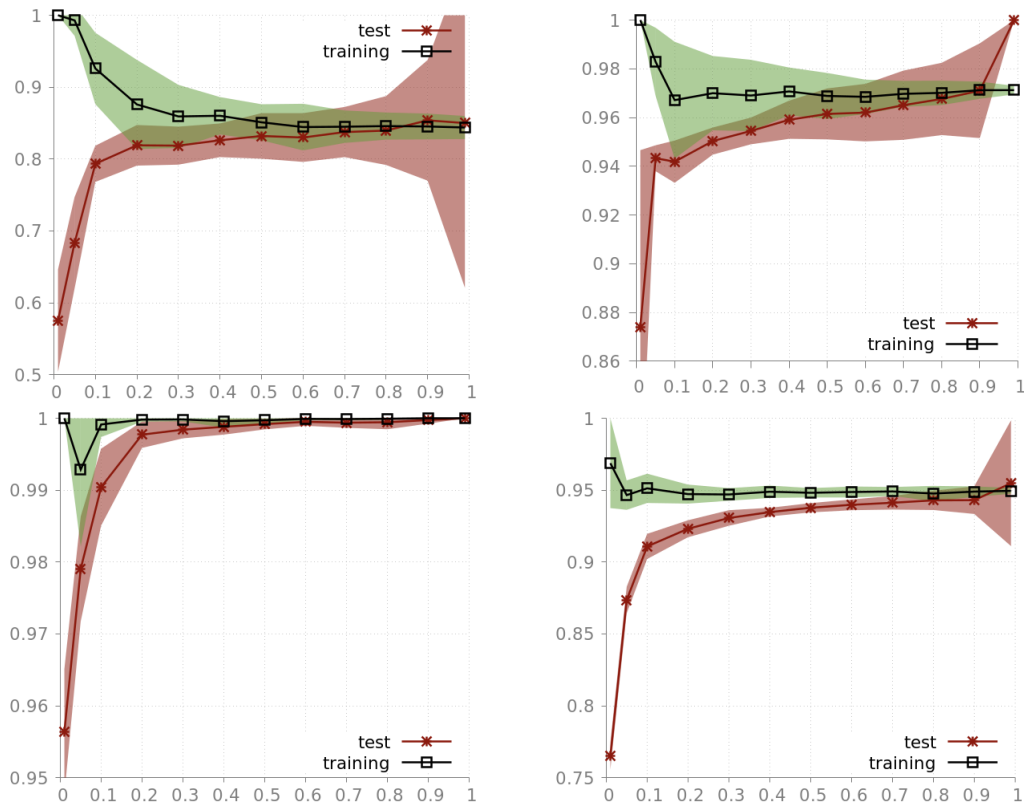


Figure 9: Learning curves for `heart`, `breast`, `mushrooms` and `splice` datasets. Despite a bias and/or irreducible error observed on `heart`, `breast` and `splice`, . The model variance appear very low on all datasets. Conversely, the bias and/or irreducible error ranges from very low on `mushrooms`, relatively low on `breast` and `splice` to high for `heart`.

The learning curves are calculated as the average over 10 runs with random




Figure 10: Learning curves on **phishing** and **skin** datasets. The accuracy on the test set is higher than on the training set. On the training set, the accuracy starts by increasing and remaining globally stable (with a drop for **skin**).

splits and are shown in Figures 11, 9 and 10. In addition, we displayed the variance of the 10 runs. The first remark concerns the variance. On all
685 datasets, the variance is very low and the accuracy on the test set extremely close if not higher than those of the training set (see **phishing** and **skin** on 10). Despite a relatively low accuracy, HCBR seems to reach a good bias-variance compromise on **heart** suggesting a more complex model space might help. In general, the remark applies on the four datasets displayed
690 in Figure 10 as the observed error rate results from bias and not only from irreducible error. Indeed, the literature comparison provided by Table 5 proves the accuracy could be improved.

The learning curves for **phishing** and **skin** are surprising for two reasons. First, the test accuracy is significantly higher than the training accuracy.
695 Two reasons might have explained this: 1) we did not use stratified sample to rebalance the dataset, so if there is a difference in prevalence between the training set and test set, the model response might be different in favor of the test set and 2) the prediction phase on those two datasets uses an heuristic such that if a case is already in the casebase, the output must be
700 similar to the one we know independently of what the model predicts (i.e. we assume there is no noise in the outputs). For 1) a comparison of the performances indicators between the test set (see Table 2) and the training set (not shown for size reason) do not show any striking difference. Regarding 2), for both datasets, the difference in accuracy between the two curves is

705 much higher than the possible gain due to the heuristic. The learning curves for the same experiment without the heuristic are display in Figure 12 and it seems to explain perfectly the phenomena. Notice that this is specific to those two datasets that contains several redundant points with the same output. For instance, the heuristic is activated also on `splice` but as the number
710 of redundant elements is very low w.r.t. to the dataset size, the impact on the accuracy is not significant. Secondly, we notice a significant drop in the training accuracy for large training set while the test accuracy remains stable. The fact the training accuracy starts relatively low is a statistical artifact due to the size of the training set as confirmed by the learning curve on very
715 small training set (see Figure 13).

Finally, the learning curve of `adult` in Figure 11 shows once again a small variance but high bias. Surprisingly, the accuracy slightly decreases with the training set size.

The analysis of the learning curves seems to indicate that the main
720 limitation of HCBR lies in a large bias. This may come either from the model space complexity or the model selection method that does not properly fit the parameters. The number of parameters in the model is equal to the cardinality of μ , itself proportional to the number of atoms in the training set. However, a closer look at how a decision is taken reveal that each case has
725 its own *small model* as a convex combination of its features. In other words, the real number of parameters to model the decision for a given case never exceed its number of features. This might not be enough to represent fairly complex functions. To overcome this problem, we can extend the model space to include additional parameters and determine a method to infer their
730 value from the training set. To discard the second hypothesis, we conducted additional investigations on the model selection in the next paragraph.

4.2.3. Model Space limitation

We are interested in understanding whether it is possible for HCBR to overfit the training set or at least improve the significantly the accuracy or MCC on the training set. We consider the matricial formulation of the support,

$$S = W\mu \tag{23}$$

and we are interested in knowing whether or not we can perturbate μ such that S leads to classify correctly all the examples. As we already verified that
735 the initial μ provides better results than baseline models, we would like to




Figure 11: Learning curve for **adult**.




Figure 12: Learning curves on **phishing** and **skin** datasets without the heuristic. Compared to Figure 10, the learning curve behave as expected in theory.

perturbate as little as possible μ . For this, we calculate \mathbf{k} such that $S + \mathbf{k}$ would correctly classify all the examples. We are looking for δ such that,

$$S + \mathbf{k} = W(\mu + \delta) \quad (24)$$

$$\Leftrightarrow S + \mathbf{k} = W\mu + W\delta \quad (25)$$

$$\Leftrightarrow \mathbf{k} = W\delta \quad (26)$$

$$\Rightarrow \delta = W^+ \mathbf{k} + [I - W^+ W] \mathbf{w}, \quad \forall \mathbf{w} \in \mathbb{R}^N \quad (27)$$

with W^+ the Moore-Penrose pseudo-inverse of W . When W is not of full rank, the solutions of the undertermined system are given for any vector




Figure 13: Learning curves on **phishing** and **adult** for very small training set size.

740 $\mathbf{w} \in \mathbb{R}^N$, however, it can be shown that $\delta = W^+ \mathbf{k}$ is a least-square minimizer, i.e. $\forall \mathbf{x} \in \mathbb{R}^M \ ||W\mathbf{x} - \mathbf{k}||_2 \geq ||W\delta - \mathbf{k}||_2$. In particular, if $||W\delta - \mathbf{k}||_2 = 0$, then all the elements would be correctly classified.

Of course, it might not be possible to obtain δ such that $||W\delta - \mathbf{k}||_2 = 0$, but it does imply that there exist no couple (\mathbf{k}, δ) such that the accuracy is 1.
 745 Solving directly for all such couples seem to be a difficult problem without additional assumptions and thus, we adopted a slightly different approach: instead of fixing \mathbf{k} a priori, we formulated the problem as optimizing the Matthew Correlation Coefficient:

$$\delta^* = \max_{\delta \in \mathbb{R}^M} \text{MCC}(\delta, W) - c||\delta||_2^2 \quad (28)$$

where MCC is the Matthew Correlation Coefficient associated to $S = W\delta$ and
 750 c a regularization factor. The regularization factor translates the idea that a smaller perturbation to obtain the same MCC is better than a larger one. We arbitrarily set c to 0.1. Despite further work could be needed to determine a more tailored value, the conclusions drawn from this section would remain valid.

755 To solve (28), we used a $(\mu + \lambda)$ genetic algorithm with an evolution strategy. Each individual is made of two vectors: $\delta \in \mathbb{R}^M$ and a $\nu \in \mathbb{R}^M$ representing the mutation probability of each corresponding component of δ . The mutation operator is a centered gaussian perturbation and the crossover a 2-points crossover. The implementation is provided by DEAP [30].

760 We performed 10 runs of a 10 folds cross-validation with random split

and compared to the result obtained without the optimization process. The dataset `mushrooms` has been discarded as HCBR reached an accuracy of 1. We set the population to 100, and for each dataset we adjusted the number of generations and the standard deviation of the gaussian mutation manually (see Table 8). To set the standard deviation of the mutation, we used $\sigma = \frac{\mu^-}{\alpha}$ where α is an empiracally determined factor and μ^- defined by $\min_{i,j} \mu_i - \mu_j$ i.e. the minimal difference between two components of μ . The rational behind is that, once again, we would like to slightly perturbate μ and it is reasonable to think that a perturbation should be small enough not to directly switch the estimation of two elements of μ . To conclude, we used a Wilcoxon signed-rank test at 5% and 1% on the test MCC obtained with and without the optimization process. There are three possible scenarios: 1) if the MCC can be improved on the training set and on the test set, it means the problem comes from the model selection method (estimation of μ and training) and results might be improved without changing the model space, 2) if the MCC can be improved on the training set but remains the same or deteriorate on the test set, the model space can represent the training set but overfit and thus should be extended 3) if the MCC cannot be improved even on the training set, then the model space is definitely not capable of representing properly the underlying decision mapping and should be extended.

A summary of the results are provided in Table 7 and the evolution of the fitness by Figure 14. A look at Figure 14 confirms that the genetic algorithm converged or was closed to converged and was able to optimize the cost function. In all cases, the optimization procedure succeed to find a vector δ that yields a better MCC on the test sets. However, the improvement is quantiatively different from a dataset to another. For instance, on `heart` the absolute difference in MCC is 16% (relative difference: 23.64%) while on `phishing` the difference is barely 1%. In general, the higher is the intial MCC and the less the improvement is visible.

The variations of MCC on the test set are mitigated. The procedure returns significant changes in only two cases (`adult`) and (`skin`) for which one is improved (`adult`) and one deteriorated (`skin`). This seems to indicate that the model selection and training phase described in Section 3.3 and 3.4 returns one of the best MCC achievable within the model space. More precisely, the vector μ represents one of the best support approximation in its neighborhood.

Table 7: Results obtained on solving (28). The columns *initial MCC*, ΔMCC training and ΔMCC test represent respectively the initial MCC on the training set, the difference of MCC obtain on the training with and without the genetic algorithm, the difference of MCC obtain on the test sets with and without the genetic algorithm. *WSR $r\%$* indicates the result of the Wilcoxon signed-rank test at $r\%$ risk (yes for significant difference in the sample median, no otherwise).

Dataset	initial MCC	ΔMCC training	ΔMCC test	WSR 5%	WSR 1%
adult	0.5190	0.0400	0.0084	yes	yes
breasts	0.9360	0.0312	-0.0025	no	no
heart	0.6912	0.1634	-0.0023	no	no
phishing	0.8690	0.0093	0.0002	no	no
skin			-0.0118	yes	yes
splice	0.8866	0.0208	0.0006	no	no

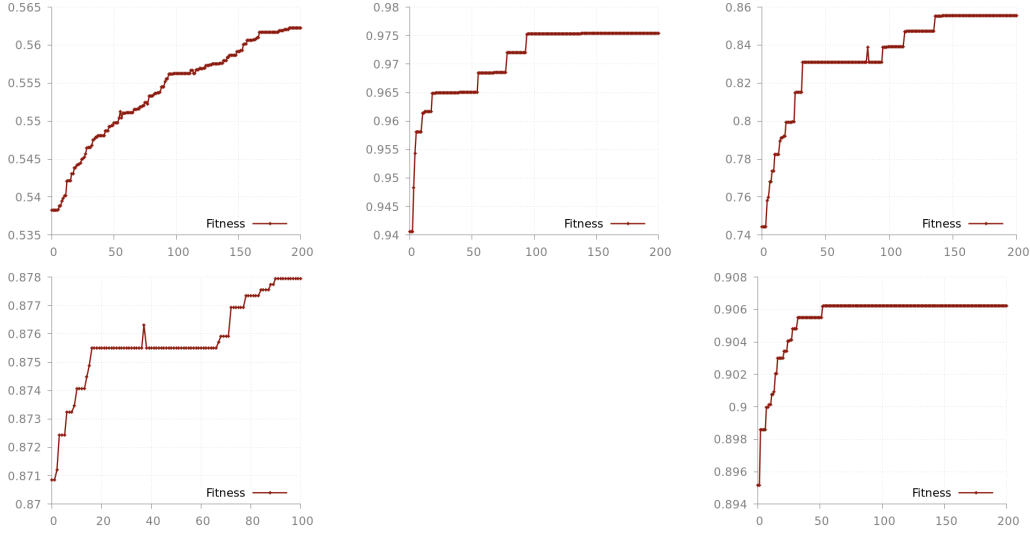


Figure 14: Maximal fitness value in the population in function of generations for **adult**, **breast**, **heart**, **phishing**, **splice** and **skin**. The fitness is lower than the MCC by definition of (28).

4.2.4. Hyperparameters η

We used a 90-10 split and set $\eta_{-1} = \eta_1$ to ease the visualization. Instead of using the decision function defined by (22), we did not produce a prediction
800 if the constraints C_1 or C_0 were not respected. It can be viewed as creating

Table 8: Hyperparameters for the genetic algorithm. *Mutation σ factor* is a factor used to set the standard deviation of the gaussian mutation. We use $\sigma = \frac{\mu^-}{\alpha}$ where α is the factor and μ^- defined by $\min_{i,j} \mu_i - \mu_j$.

Dataset	Generations	Mutation σ factor
adult	200	10000
breasts	200	10
heart	200	10
phishing	100	1000
skin	100	1000
splice	200	100

a third class *unknown* for which we consider we cannot produce a decision. We measured the accuracy and the test set size ratio for which a prediction has been produced for different values of $\eta := \eta_{-1} = \eta_1$. If \bar{J} correctly approximates J , increasing η should increase the accuracy while the test set ratio should remain high. Additionally, we plot the test set ratio in function of the accuracy and calculate the Pareto frontier⁸ which represents the best compromises accuracy/ratio. The closer the points are to (1, 1) the better it is. A Pareto frontier consisting of (1, 1) represents the perfect model (e.g. reached on **mushroom** dataset). Figures 15, 16, 17 and 18 provides the result for the best and worst two datasets. Figure 19 shows all of the four Pareto frontiers. As expected, the results are better on **phishing** and **breast**. On **phishing**, **breast** and **heart**, the accuracy globally increases with η while on **heart** the accuracy slightly decreases indicating poor influence of the hyperparameters and model.

Notice that for certain values of η it is possible to reach 100% accuracy with **heart** while it is not with **breast**. Also, for high values of η , we observe a fall in accuracy for **breast**. We suspect those two phenomena to appear because we used the same value for η_0 and η_1 .

More work is required to fully study the influence of the hyperparameters ($\eta_{-1}, \eta_1, \bar{\eta}_{-1}, \bar{\eta}_1$) and how to select l_{-1} and l_1 . We believe it is the key to improve the overall performance, and it is possible to derivate the best values from the training set. Also, a comparison with binary classification methods

⁸Points such that improving one component would degrades the other one.

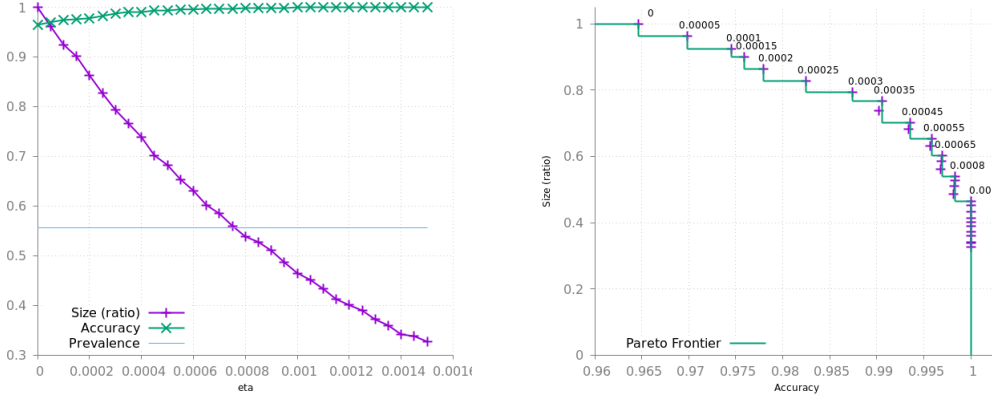


Figure 15: Influence of η on **phishing** dataset.

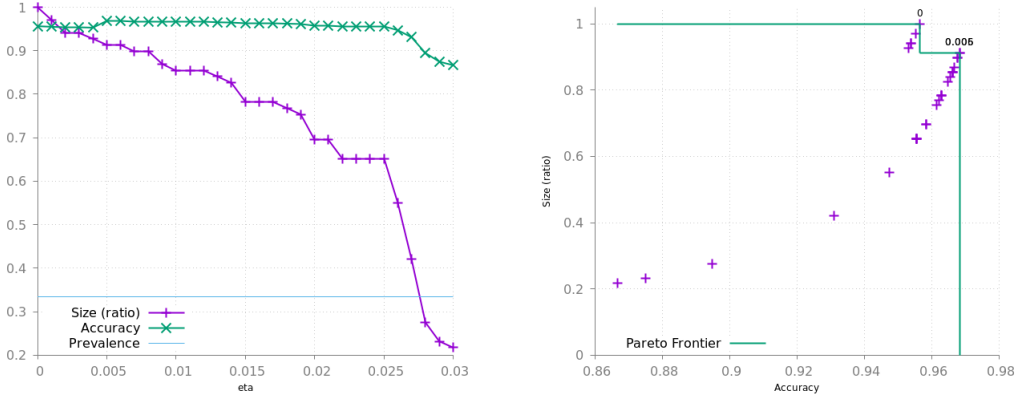


Figure 16: Influence of η on **breast** dataset.

that provide a prediction confidence metric is necessary.

5. Conclusion

825 This paper presented HCBR, a method for binary classification in unstruc-
 tured space using a hypergraph representation of information and building a
 convex combination out of the induced partition to determine the support
 for each class. The general framework introduced by (10) is instantiated in
 Section 3.3 and 3.4 where the support is determined using all the interactions
 830 between the hyperedges. Beyond this specific implementation, one can imag-
 ine different model selection methods to be used, e.g. using some assumptions
 on the data.

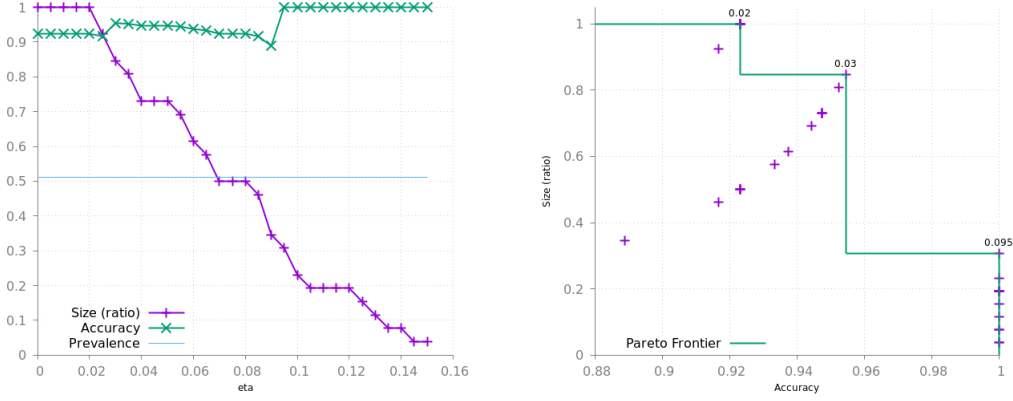


Figure 17: Influence of η on **heart** dataset.

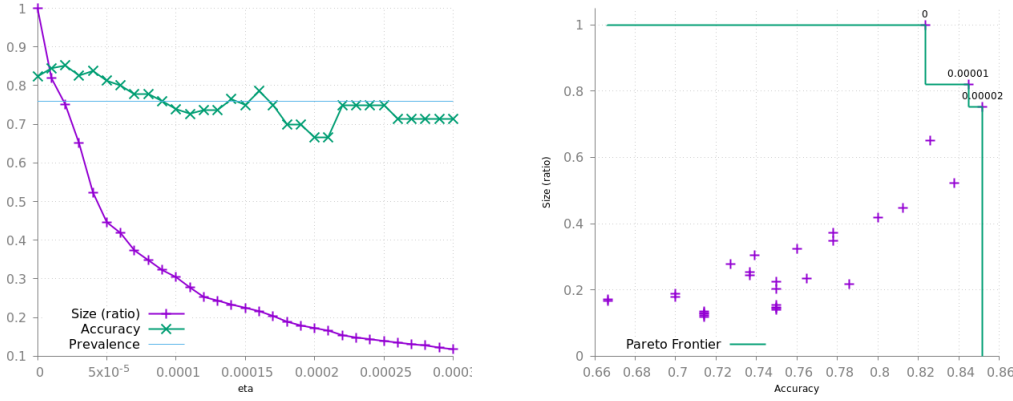


Figure 18: Influence of η on **adult** dataset.

However, being totally agnostic on the data representation is convenient compared to many methods such as SVM. It allows combining information from multiple sources by simply *stacking* the information. It does not require transforming the data to fit the algorithm, often by designing specific ad-hoc metrics.

HCBR has been tested on 7 well-known datasets and demonstrated similar accuracies when compared to the best results from the literature. Additionally, we performed a comparison with 9 alternative methods to find out HCBR along with Neural Network outperform in average with a constant good result. This result is important w.r.t. to the capacity to HCBR to be easily used and deployed in practice. We showed that the model is properly




Figure 19: Pareto Frontiers comparison.

calibrated i.e. the support measure can be seen as a confidence measure on the prediction. We empirically validated the exponential worst-case complexity. Finally, we studied the properties and limitations of the model space. The learning curves showed that the model does not overfit on any dataset, has a small variance but tend to have a large bias. To determine if the problem comes from the model space complexity or an inefficient model selection procedure, we used a genetic algorithm to optimize an auxiliary function and have shown that the model selection procedure provides one of the best possible performance within the model space. Hence, further work will focus on extending the model space.

This proof of concept raises many questions and offer many improvement axes. First of all, it seems relatively easy to extend the method to several classes but it needs an additional empirical validation. As most of the computational effort is on calculating the class support, adding more classes will linearly increase the computation time and thus, working on a faster algorithm or an approximation of the main measure should be investigated. The solution may come from exploring the feature selection capacity of HCBR. Indeed, by the hypergraph construction, it may be possible to remove from the partition some elements that do not participate enough (e.g. not being in enough cases at the same time), reducing the computation time. Additionally, we plan to investigate how to generate an explanation about each prediction and one about the decision function J itself, using not only the convex combination and the strength of the partition elements, but also the link between cases in a similar way a lawyer may use past cases to make analogies

or counter-examples. We also work on an online version of HCBR where the hypergraph is constructed case after case, including forgetting some old past cases (which would allow handling non-stationary environment). It seems also possible not only to add new examples dynamically, but also some vertices (i.e. adding some pertinent information to some cases) without generating the whole model from scratch.

Empirically, further experiments should focus on more unstructured datasets (for instance for text classification). As stated previously, strategies for hyperparameter tuning are also a priority. Last but not least, we would like to answer some questions: can we provide some quality bounds depending on the initial hypergraph and thus the dataset? How to handle continuous values without discretization?

Acknowledgment

The author warmly thanks Dr. Jean-François Puget, IBM Analytics, for their useful suggestions and advice to improve this paper.

References

- [1] A. Quemy, Binary classification with hypergraph case-based reasoning, in: Proc. of Workshop on Design, Optimization, Languages and Analytical Processing of Big Data (DOLAP), 2018.
- [2] S. Lifan, G. Yue, Z. Xibin, W. Hai, G. Ming, S. Jianguang, Vertex-weighted hypergraph learning for multi-view object classification, in: Proc. of Int. Joint Conf. Artificial Intelligence (IJCAI), 2017, pp. 2779–2785.
- [3] Y. Huang, Q. Liu, S. Zhang, D. N. Metaxas, Image retrieval via probabilistic hypergraph ranking, in: Proc. IEEE Comput. Soc. Conf. Comput. Vis. Pattern Recognit., 2010, pp. 3376–3383.
- [4] D. Zhou, J. Huang, B. Schölkopf, Learning with hypergraphs: Clustering, classification, and embedding, in: B. Schölkopf, J. C. Platt, T. Hoffman (Eds.), Adv. Neural. Inf. Process. Syst. (NIPS), MIT Press, 2007, pp. 1601–1608.
- [5] C. Berge, Hypergraphs: combinatorics of finite sets, Vol. 45, Elsevier, 1984.

- 900 [6] R. Paige, R. E. Tarjan, Three partition refinement algorithms, *SIAM J. Comput.* 16 (6) (1987) 973–989.
- [7] H. He, E. A. Garcia, Learning from imbalanced data, *IEEE Transactions on Knowledge and Data Engineering* 21 (9) (2009) 1263–1284.
- 905 [8] R. P. Datta, S. Saha, Applying rule-based classification techniques to medical databases: an empirical study, *Int. J. of Business Int. and Syst. Engineering* 1 (1) (2016) 32–48.
- [9] L. Jiang, Learning instance weighted naive bayes from labeled and unlabeled data, *J. Intell. Inf. Syst.* 38 (1) (2012) 257–268.
- 910 [10] A. M. Sagir, S. Sathasivam, A hybridised intelligent technique for the diagnosis of medical diseases., *Pertanika Journal of Science & Technology* 25 (2).
- [11] S. Das, Filters, wrappers and a boosting-based hybrid for feature selection, in: *Proc. of Int. Conf. on Machine Learning, ICML '01*, Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 2001, pp. 74–81.
- 915 [12] W. Hadi, G. Issa, A. Ishtaiwi, Acprism: Associative classification based on prism algorithm, *Information Sciences* 417 (Supplement C) (2017) 287 – 300.
- [13] F. Thabtah, R. M. Mohammad, L. McCluskey, A dynamic self-structuring neural network model to combat phishing, in: *Int. Joint Conf. on Neural Networks (IJCNN)*, 2016, pp. 4221–4226.
- 920 [14] M. A. U. H. Tahir, S. Asghar, A. Zafar, S. Gillani, A hybrid model to detect phishing-sites using supervised learning algorithms, in: *Int. Conf. on Computational Science and Computational Intelligence (CSCI)*, 2016, pp. 1126–1133.
- 925 [15] G. Kou, Y. Lu, Y. Peng, Y. Shi, Evaluation of classification algorithms using mcdm and rank correlation, *Int. J. of Information Technology & Decision Making* 11 (01) (2012) 197–225.
- [16] Y.-J. Lee, O. Mangasarian, Ssvm: A smooth support vector machine for classification, *Computational Optimization and Applications* 20 (1) (2001) 5–22.

- 930 [17] R. B. Bhatt, G. Sharma, A. Dhall, S. Chaudhury, Efficient skin region segmentation using low complexity fuzzy decision tree model, in: IEEE India Conference, 2009, pp. 1–4.
- [18] F. Ö. Çatak, Classification with boosting of extreme learning machine over arbitrarily partitioned data, *Soft Computing* 21 (9) (2017) 2269–2281.
- 935 [19] H.-L. Chen, B. Yang, J. Liu, D.-Y. Liu, A support vector machine classifier with rough set-based feature selection for breast cancer diagnosis, *Expert Systems with Applications* 38 (7) (2011) 9014 – 9022.
- [20] K. Polat, S. Güneş, Breast cancer diagnosis using least square support vector machine, *Digital Signal Processing* 17 (4) (2007) 694 – 701.
- 940 [21] M. F. Akay, Support vector machines combined with feature selection for breast cancer diagnosis, *Expert Systems with Applications* 36 (2) (2009) 3240–3247.
- [22] A. Marcano-Cedeo, J. Quintanilla-Domnguez, D. Andina, Wbcd breast cancer database classification applying artificial metaplasticity neural network, *Expert Systems with Applications* 38 (8) (2011) 9573 – 9579.
- 945 [23] E. D. Übeyli, Implementing automated diagnostic systems for breast cancer detection, *Expert Systems with Applications* 33 (4) (2007) 1054–1062.
- [24] A. Fallahi, S. Jafari, An expert system for detection of breast cancer using data preprocessing and bayesian network, *Int. J. of Advanced Science and Technology* 34 (2011) 65–70.
- 950 [25] J. R. Quinlan, Improved use of continuous attributes in c4. 5, *J. of Artificial Intelligence Research* 4 (1996) 77–90.
- [26] M. T. Cazzolato, M. X. Ribeiro, A statistical decision tree algorithm for medical data stream mining, in: *Proc. of IEEE Int. Symposium on Computer-Based Medical Systems*, 2013, pp. 389–392.
- 955 [27] S. Basterrech, A. Mesa, N.-T. Dinh, Generalized linear models applied for skin identification in image processing, in: *Intelligent Data Analysis and Applications*, Springer, 2015, pp. 97–107.
- 960

- 965 [28] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, E. Duchesnay, Scikit-learn: Machine learning in Python, *J. of Machine Learning Research (JMLR)* 12 (2011) 2825–2830.
- [29] D. Chicco, Ten quick tips for machine learning in computational biology, *BioData Mining* 10 (1).
- 970 [30] F.-A. Fortin, F.-M. De Rainville, M.-A. Gardner, M. Parizeau, C. Gagné, DEAP: Evolutionary algorithms made easy, *J. of Machine Learning Research (JMLR)* 13 (2012) 2171–2175.